

# Midterm 1 exam

---

**In-class midterm.** 2:00-2:15pm on Monday Feb 25.

- No makeups.

## Rules.

- Closed book, closed note.
- Covers all material thru Lecture 8.
- No computers or other computational devices.
- 8.5-by-11 cheatsheet (one side, in your own handwriting).

including associated  
assignments  
(but no serious Java programming)



## Midterm preparation.

- Review session: Go to recitations.

# Binary heap: practical improvements

**Caching.** Binary heap is not cache friendly.

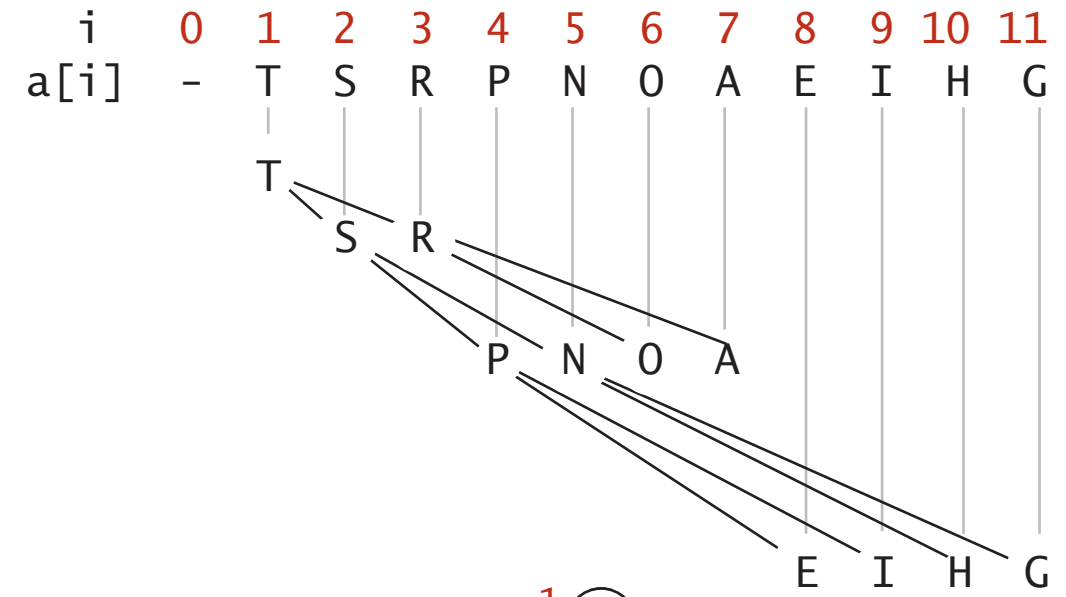
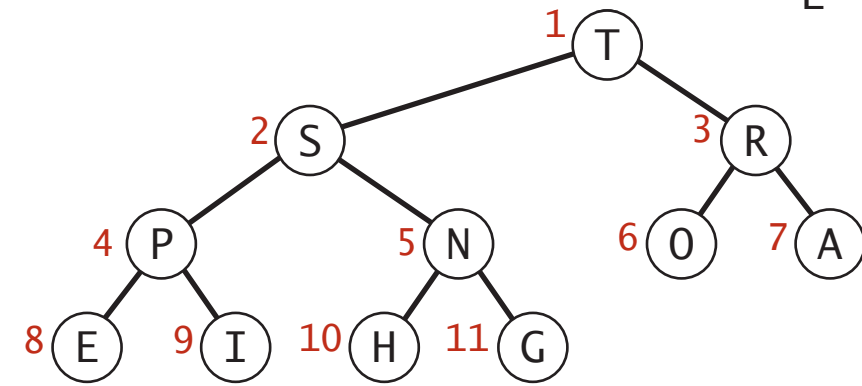
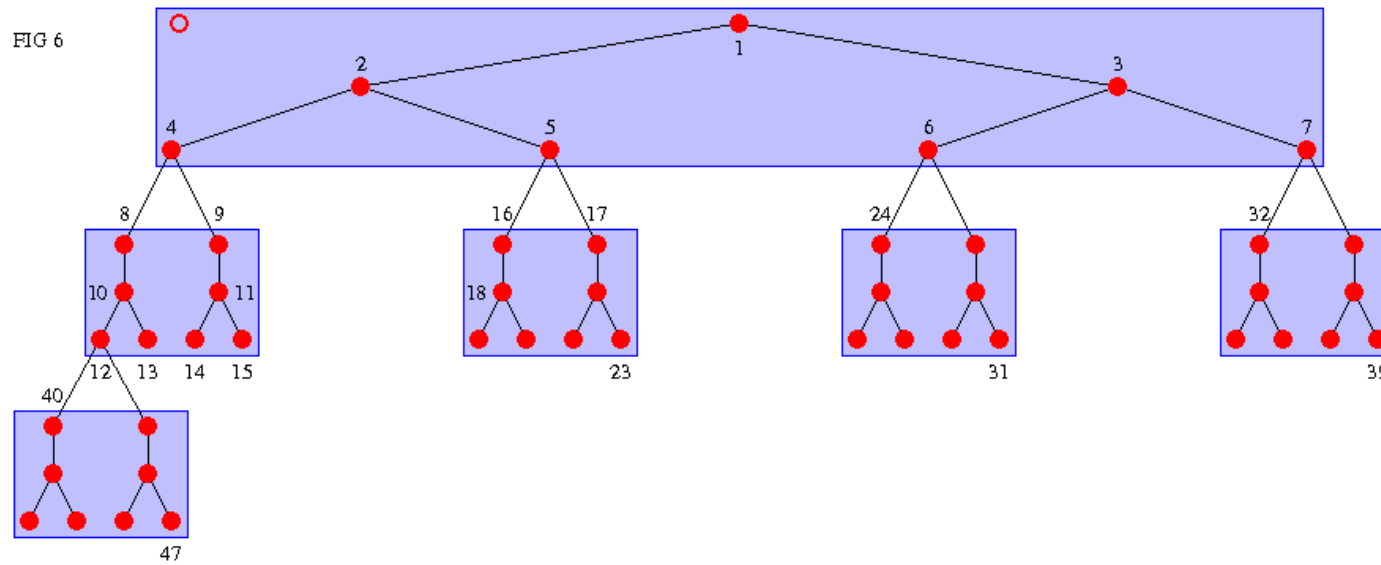
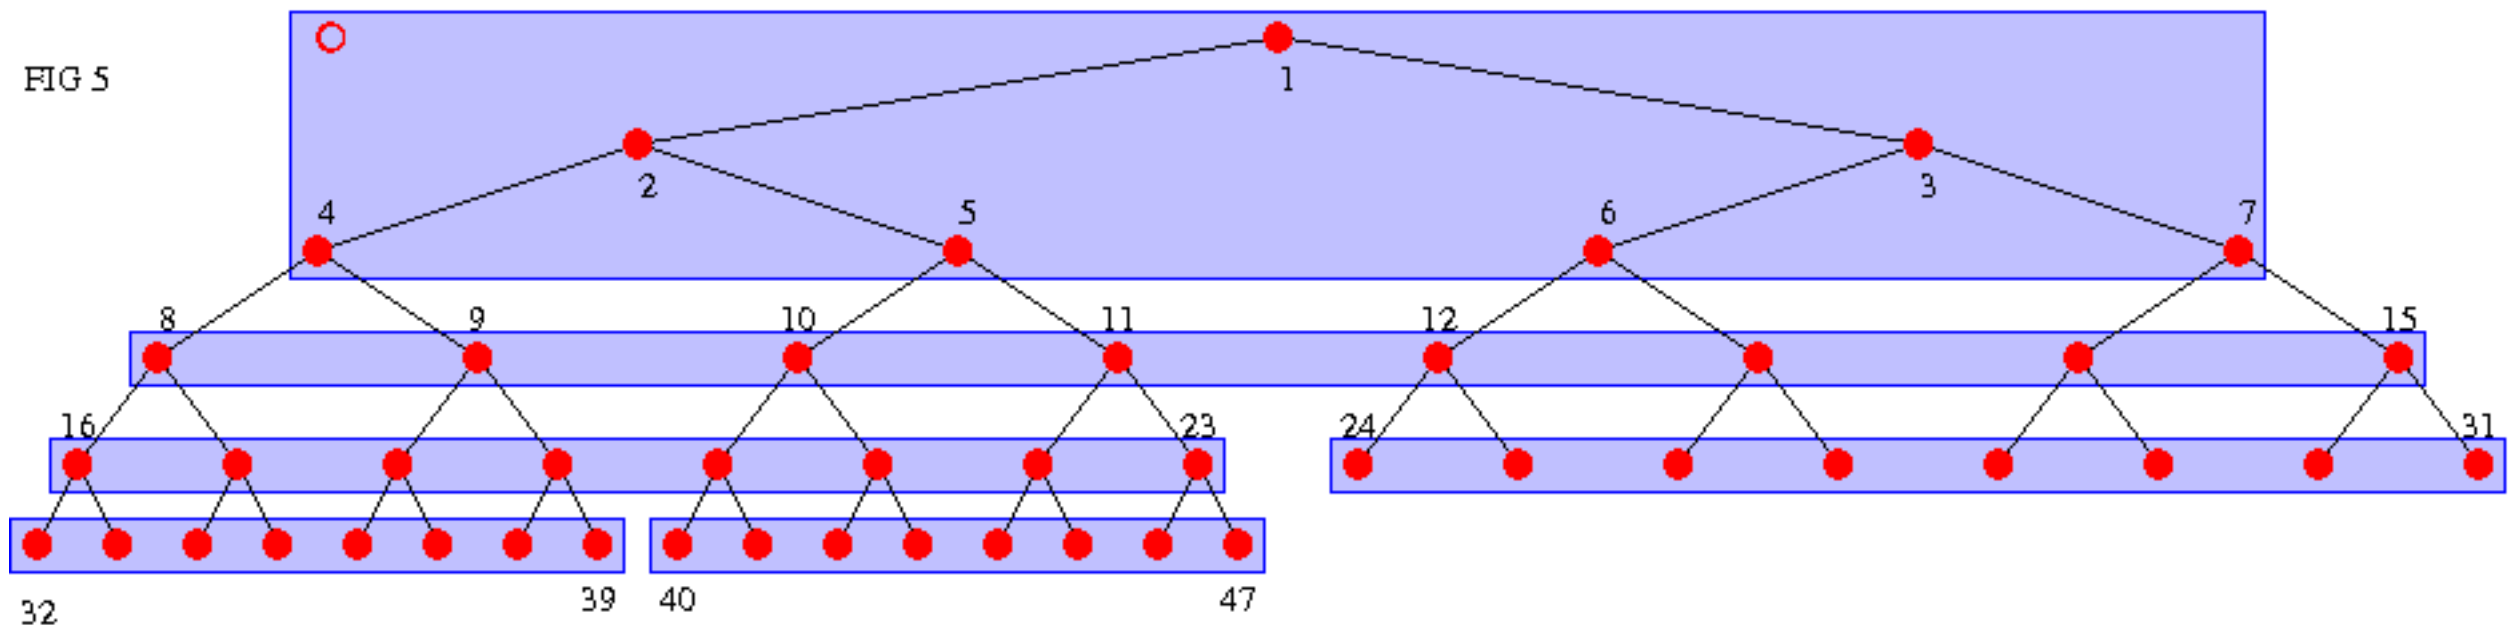


FIG 6



Heap representations

FIG 5



# 2-3 TREES

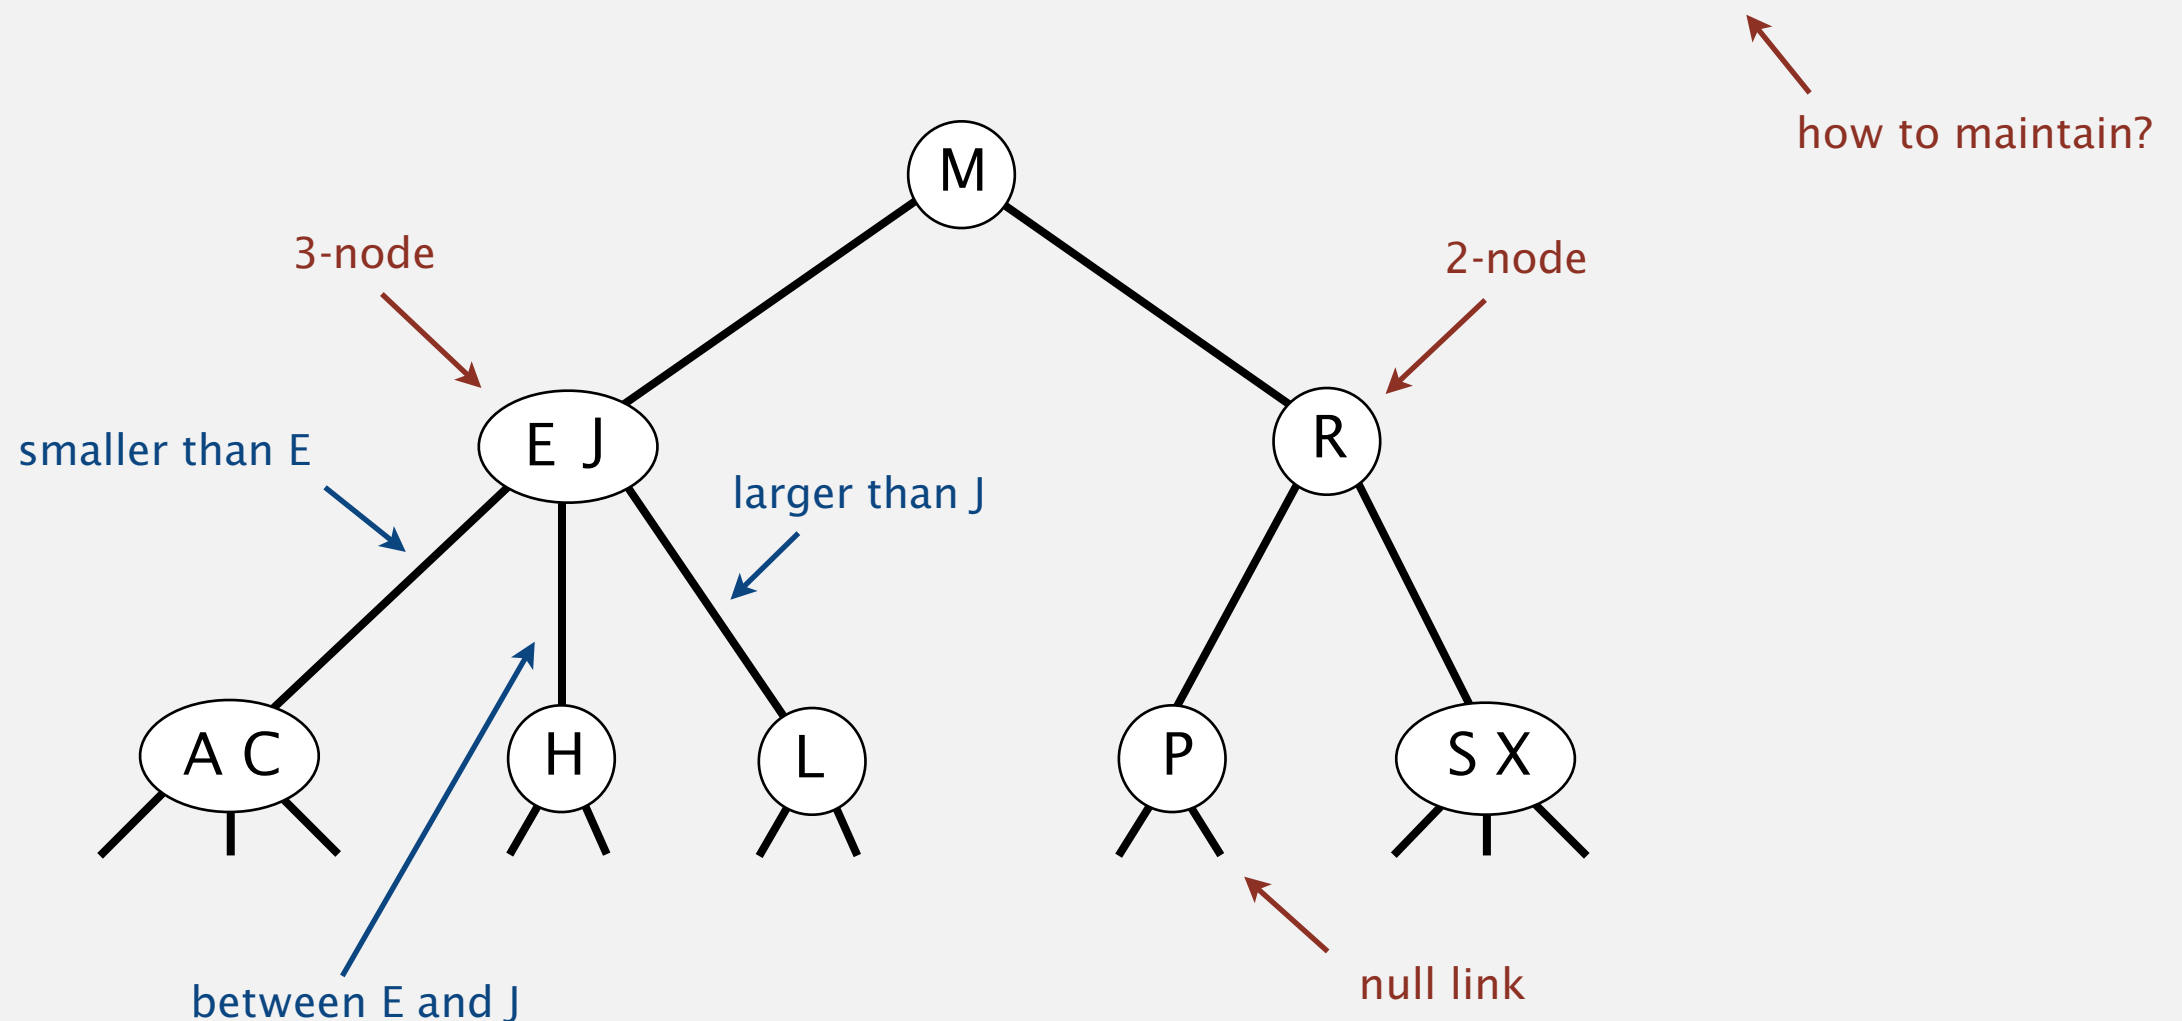
# 2-3 tree

Allow 1 or 2 keys per node.

- 2-node: one key, two children.
- 3-node: two keys, three children.

**Symmetric order.** Inorder traversal yields keys in ascending order.

**Perfect balance.** Every path from root to null link has same length.



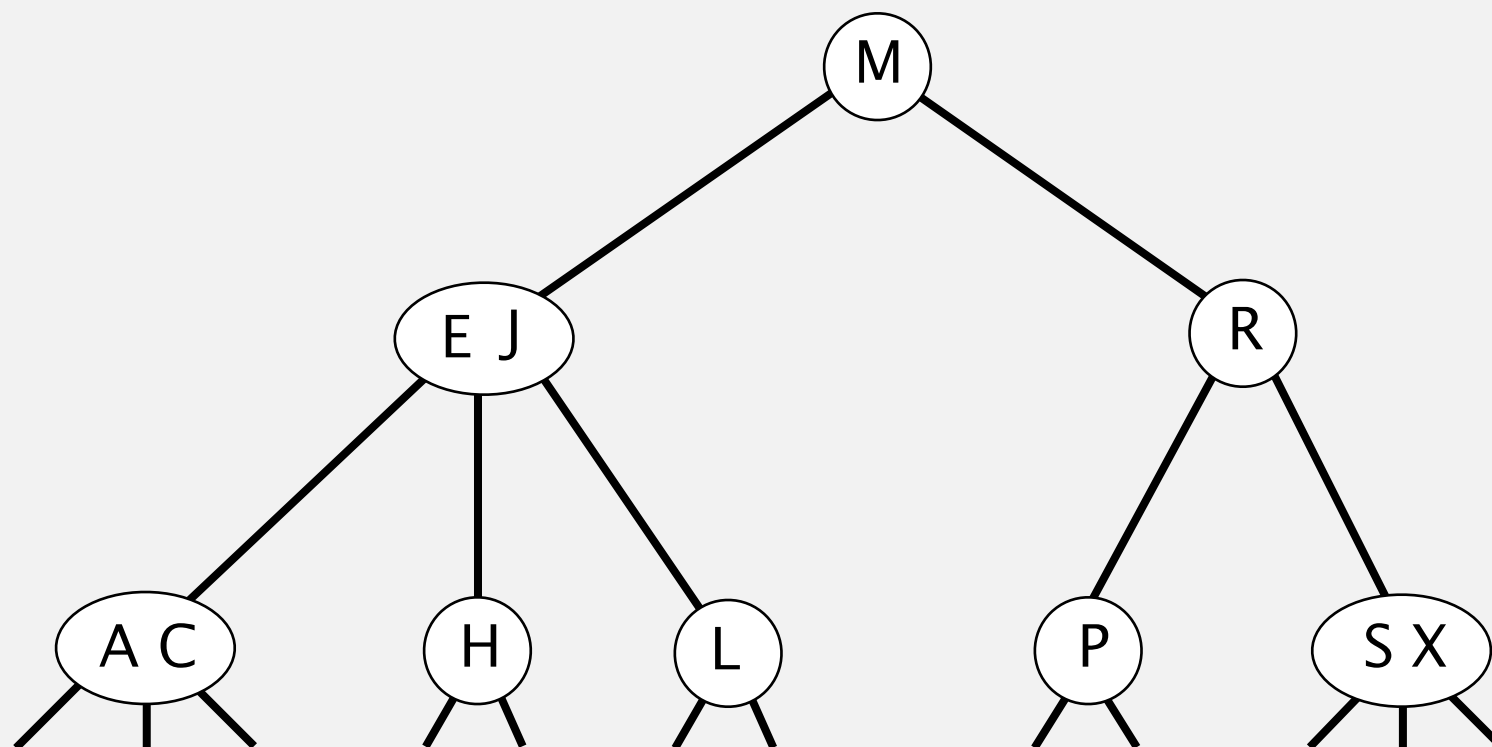
## 2-3 tree demo

---

### Search.

- Compare search key against keys in node.
- Find interval containing search key.
- Follow associated link (recursively).

search for H



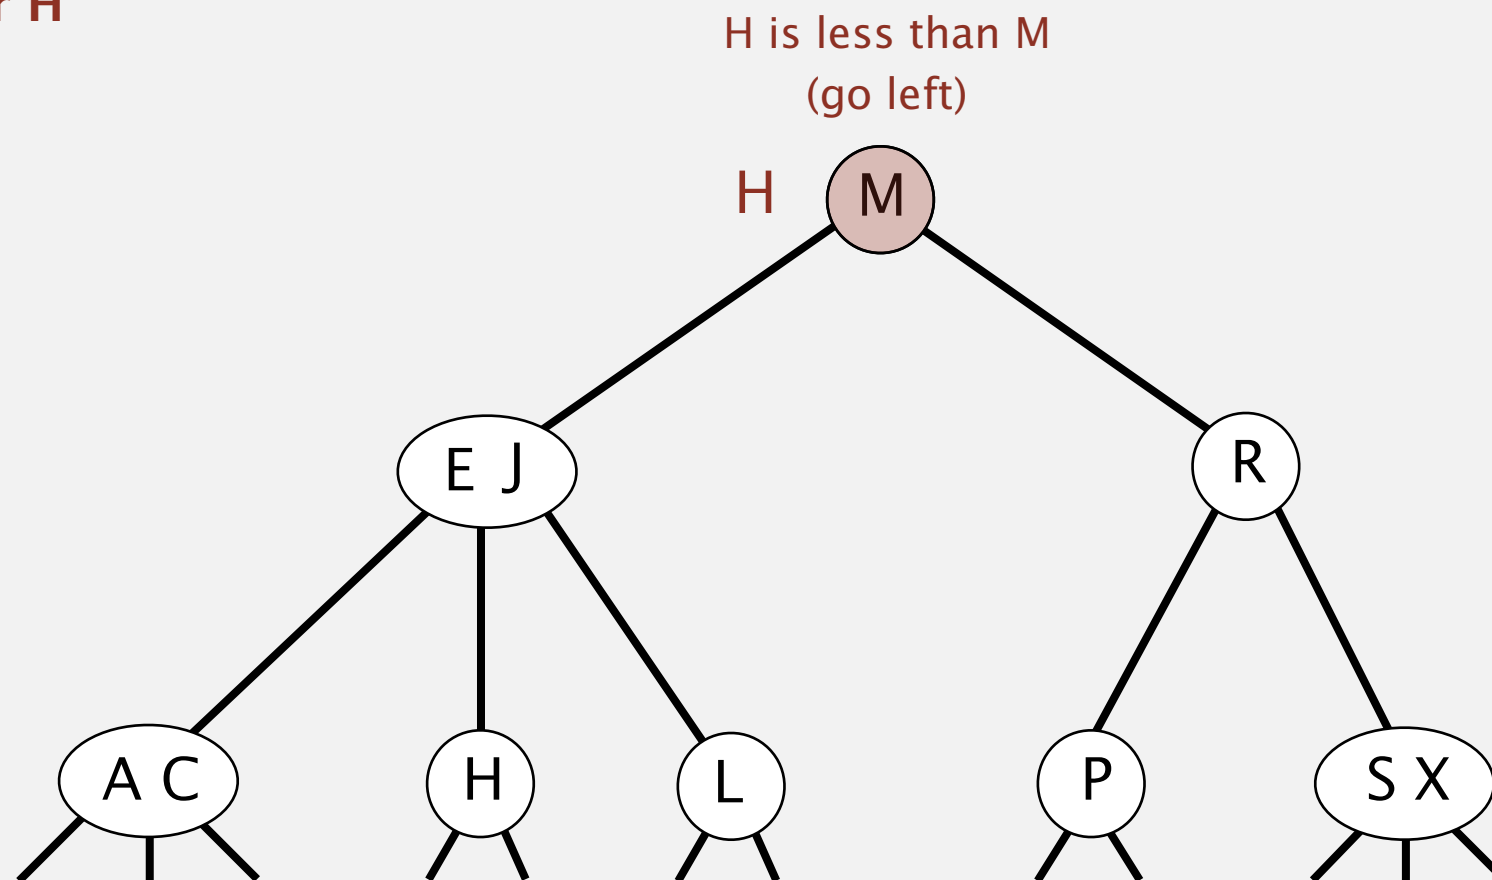
## 2-3 tree demo: search

---

### Search.

- Compare search key against keys in node.
- Find interval containing search key.
- Follow associated link (recursively).

search for H



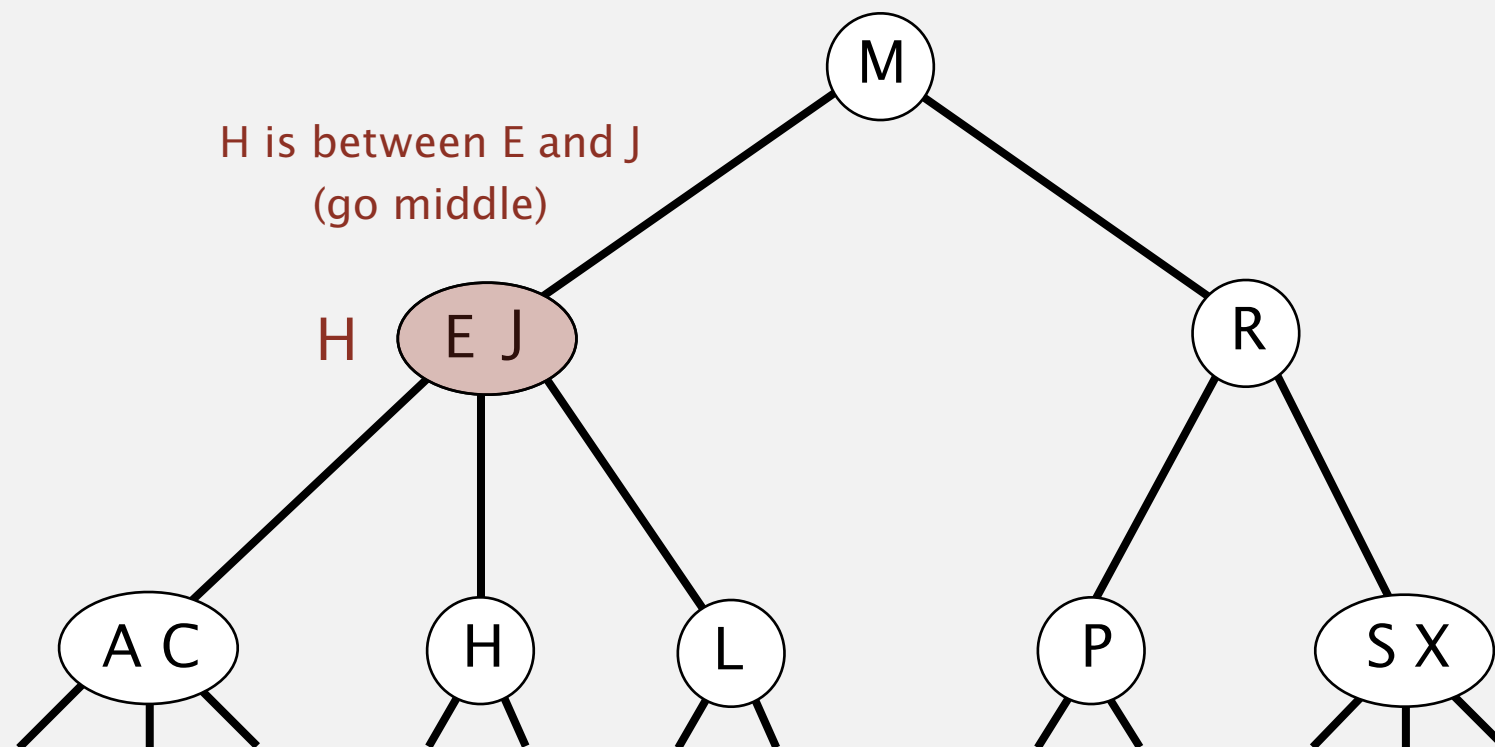
## 2-3 tree demo: search

---

### Search.

- Compare search key against keys in node.
- Find interval containing search key.
- Follow associated link (recursively).

search for H



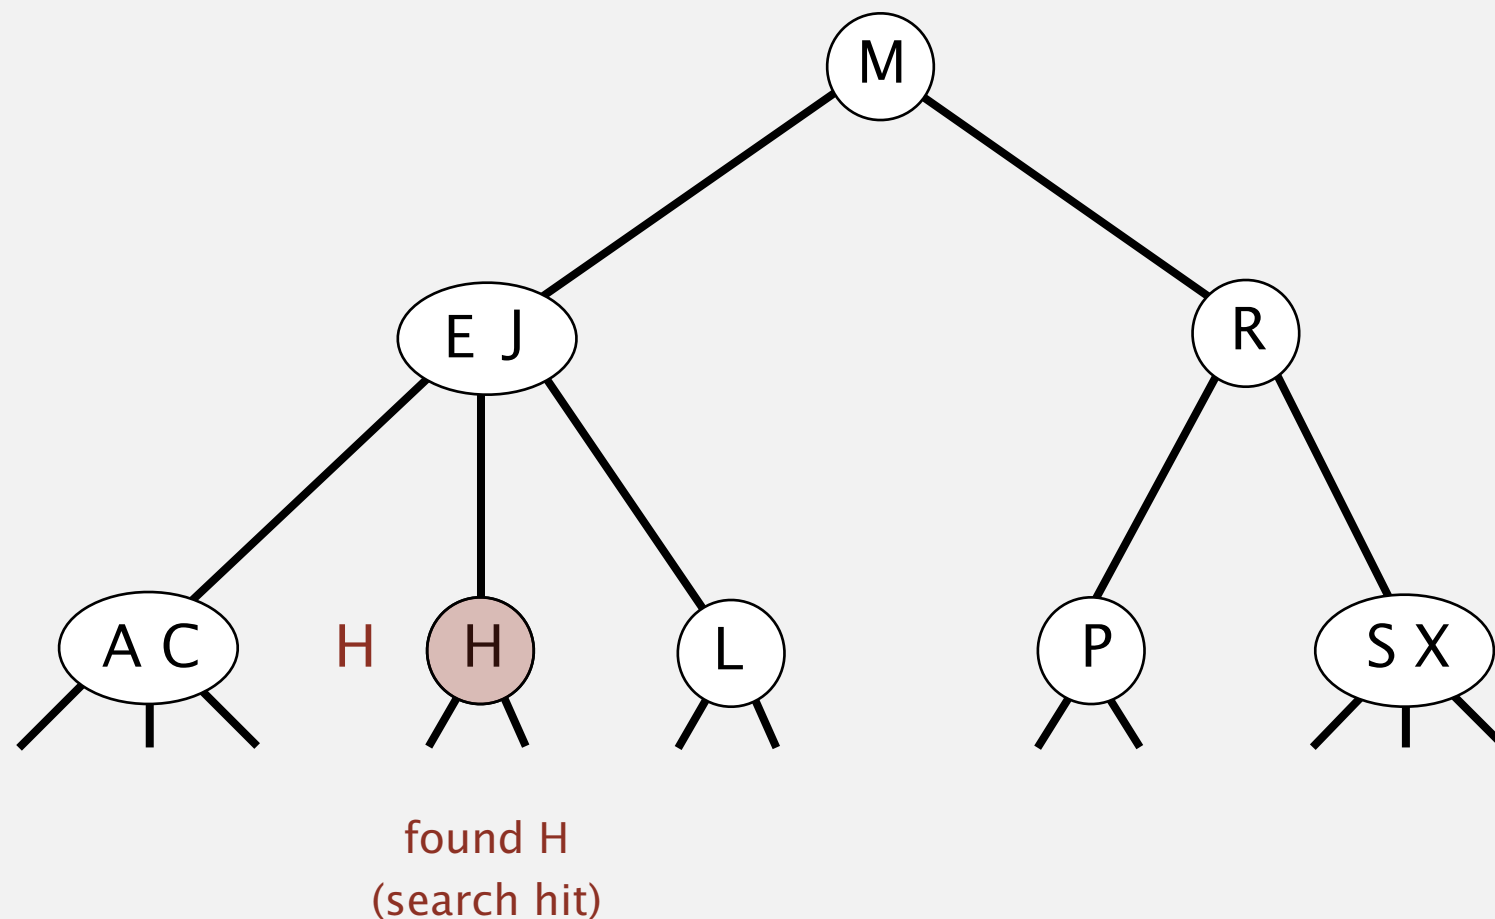
## 2-3 tree demo: search

---

### Search.

- Compare search key against keys in node.
- Find interval containing search key.
- Follow associated link (recursively).

search for H





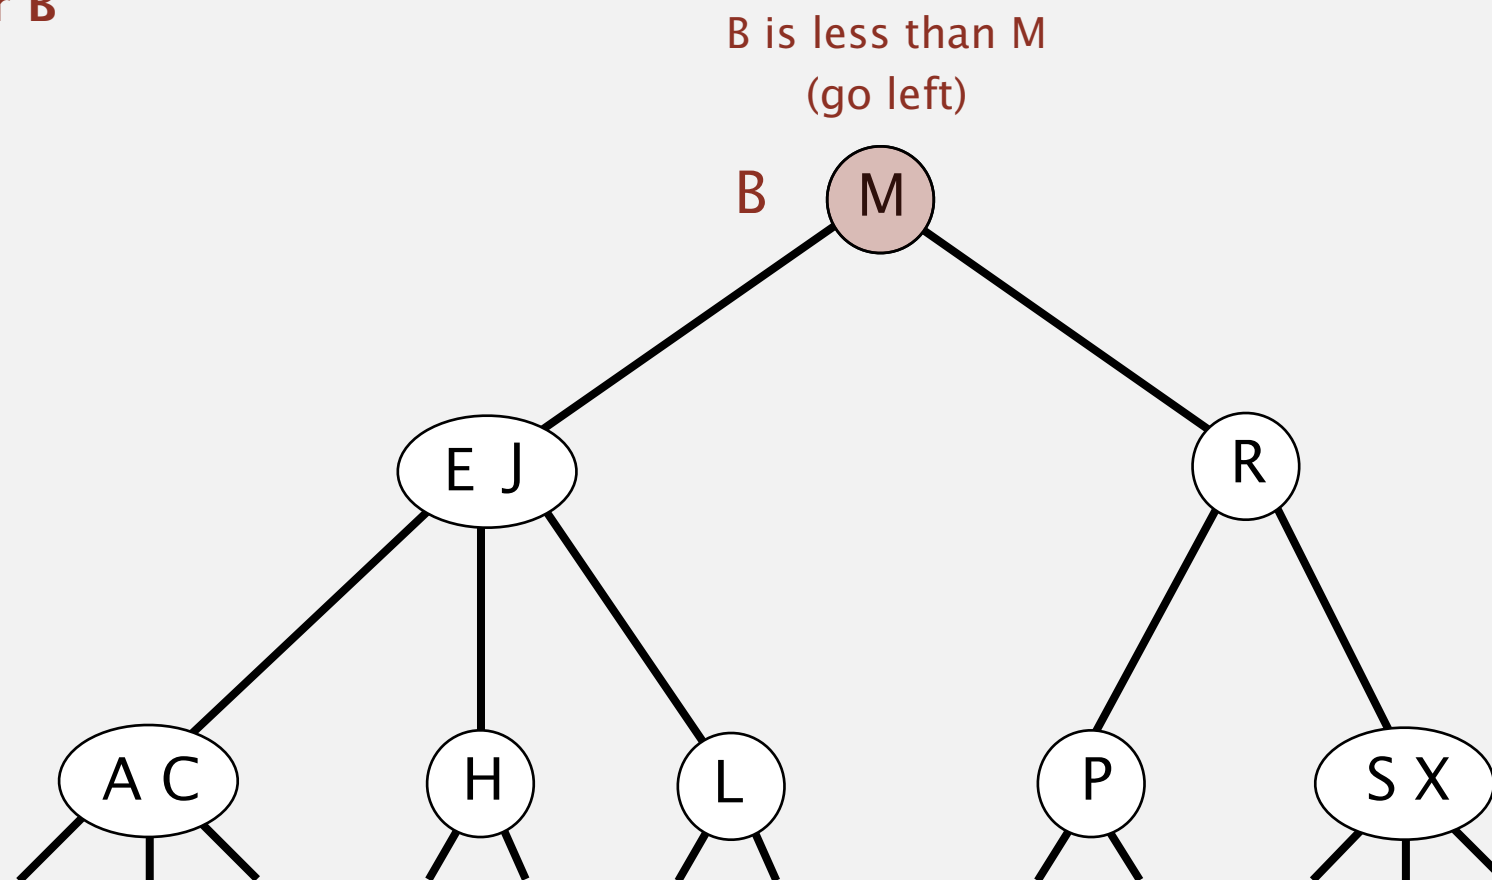
## 2-3 tree demo: search

---

### Search.

- Compare search key against keys in node.
- Find interval containing search key.
- Follow associated link (recursively).

search for B



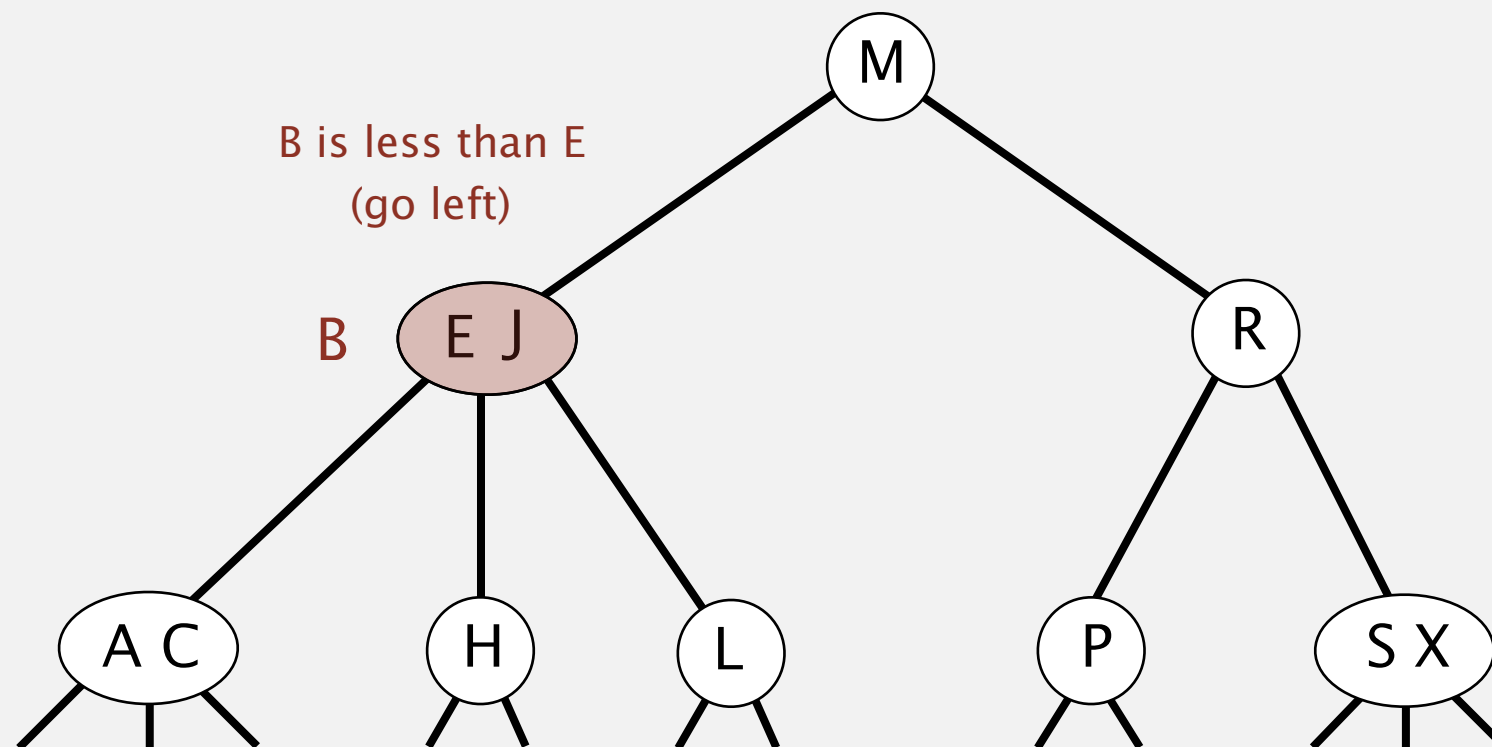
## 2-3 tree demo: search

---

### Search.

- Compare search key against keys in node.
- Find interval containing search key.
- Follow associated link (recursively).

search for B



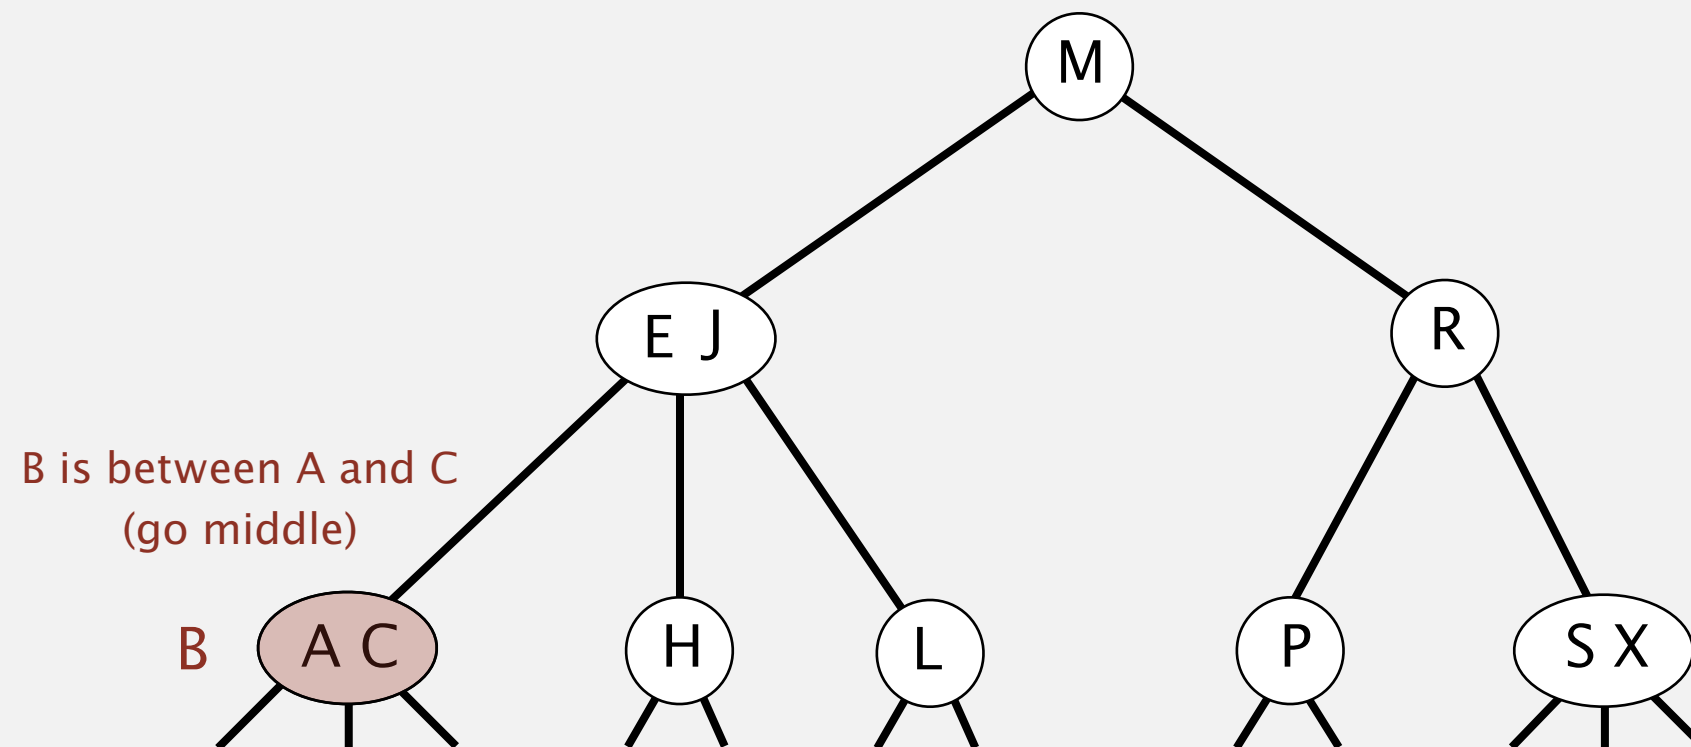
## 2-3 tree demo: search

---

### Search.

- Compare search key against keys in node.
- Find interval containing search key.
- Follow associated link (recursively).

search for B



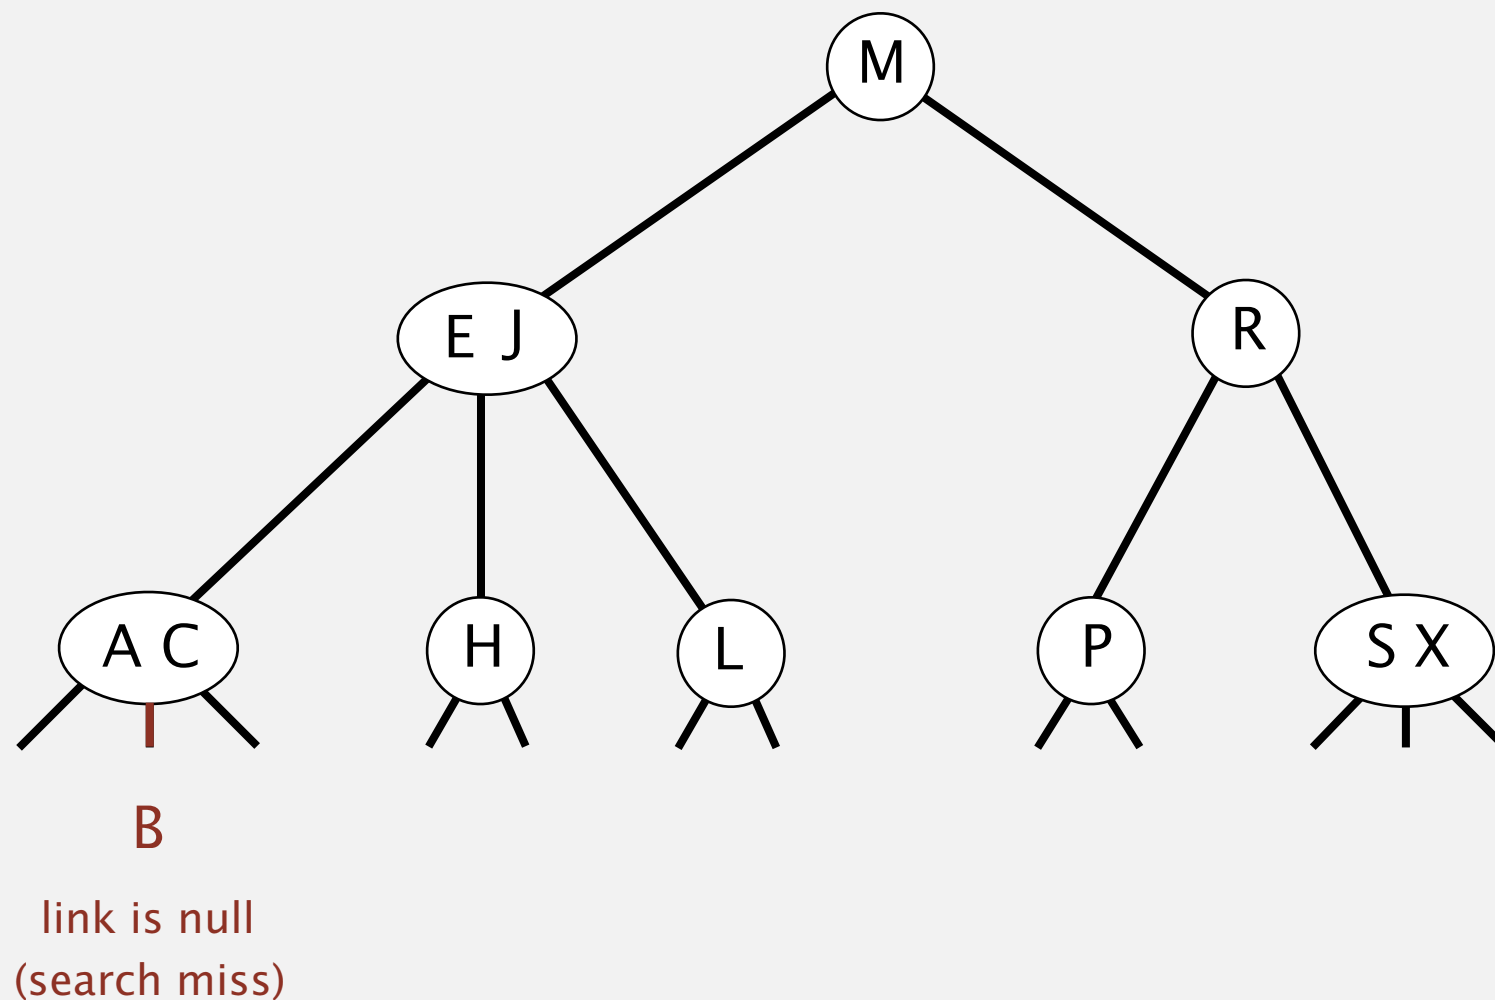
## 2-3 tree demo: search

---

### Search.

- Compare search key against keys in node.
- Find interval containing search key.
- Follow associated link (recursively).

search for B



# 2-3 TREES

Inserting Values

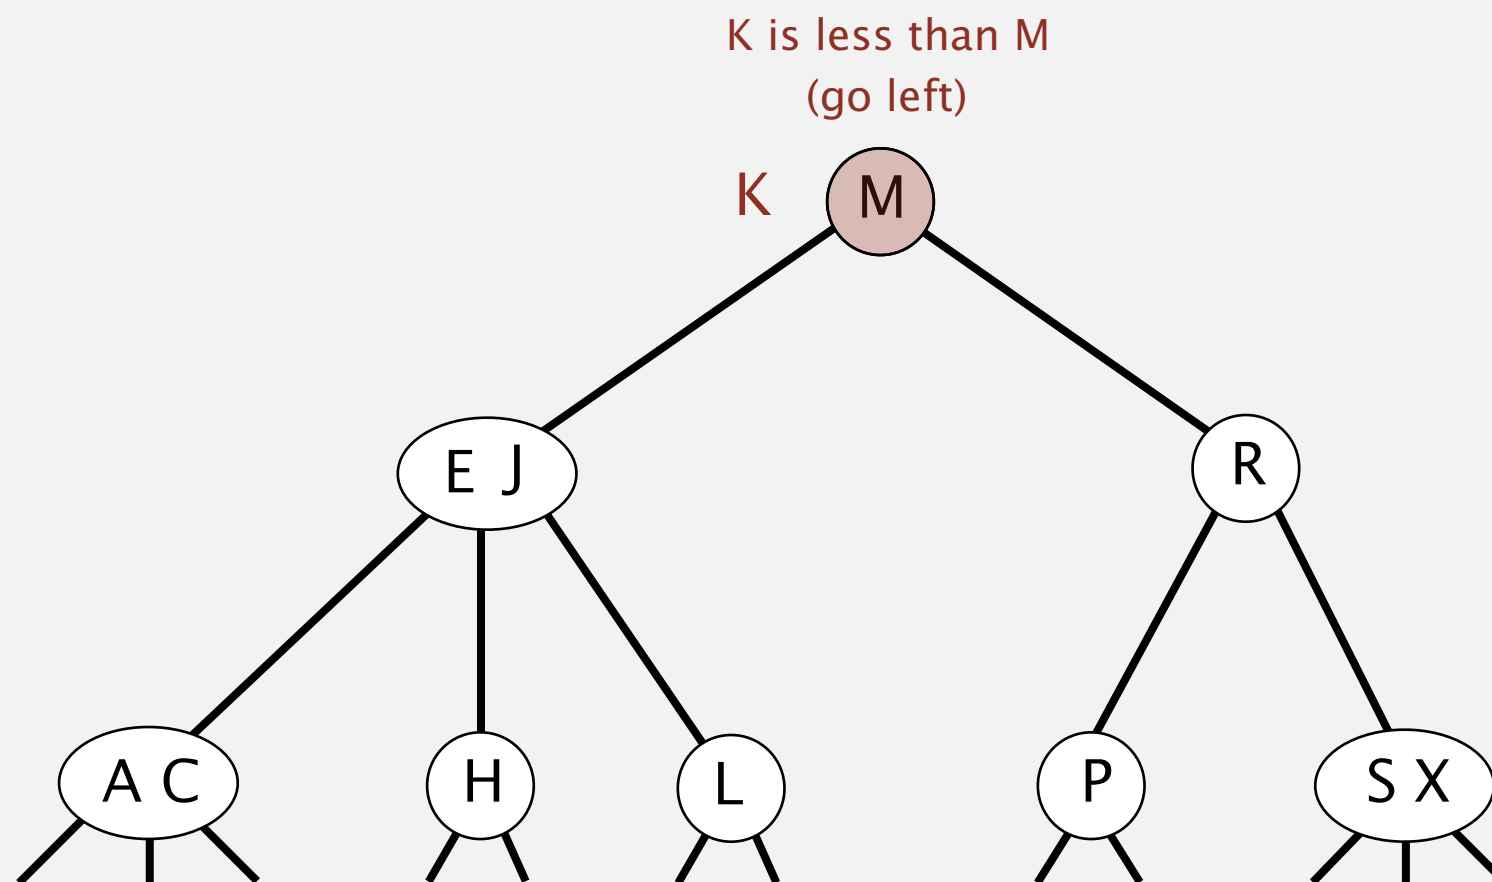
## 2-3 tree demo: insertion

---

Insert into a 2-node at bottom.

- Search for key, as usual.
- Replace 2-node with 3-node.

insert K



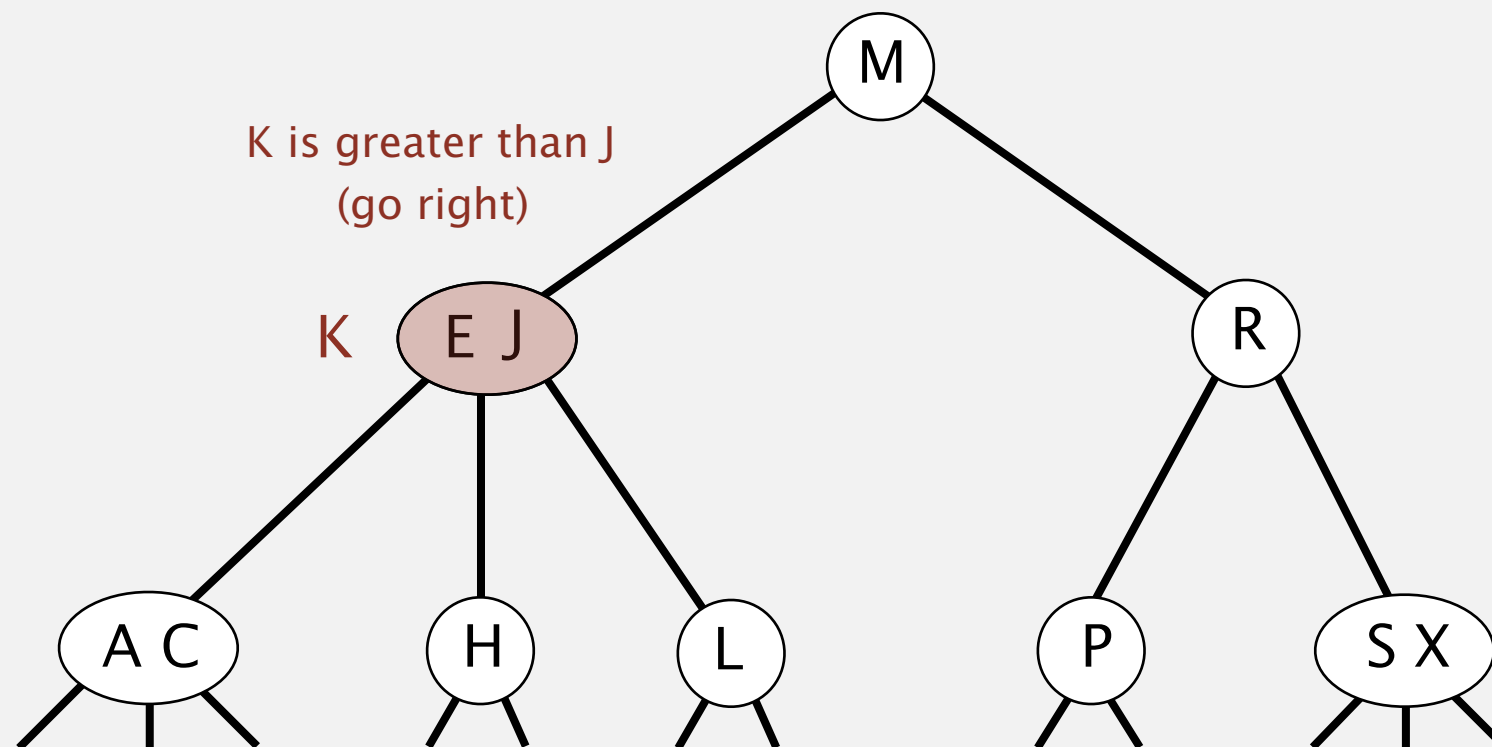
## 2-3 tree demo: insertion

---

Insert into a 2-node at bottom.

- Search for key, as usual.
- Replace 2-node with 3-node.

insert K



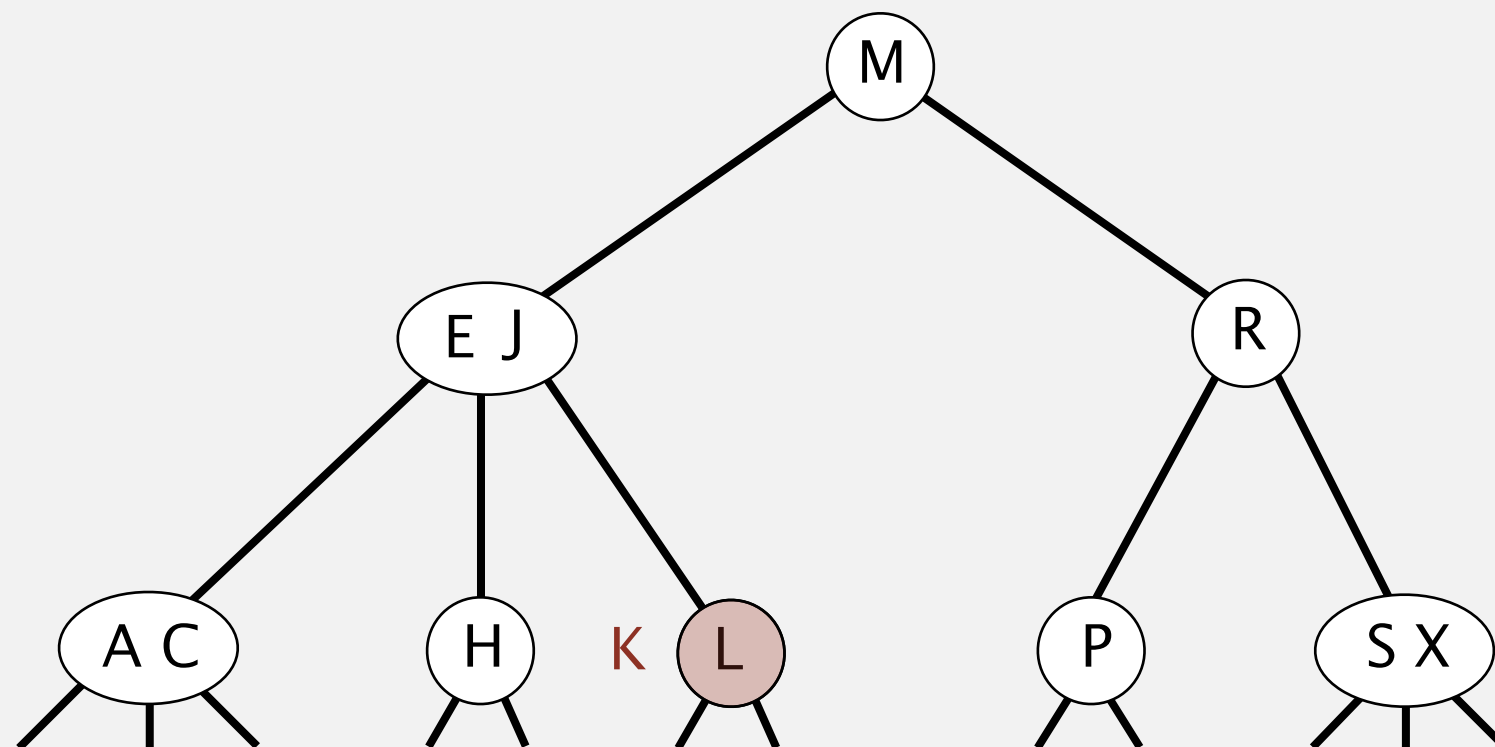
## 2-3 tree demo: insertion

---

Insert into a 2-node at bottom.

- Search for key, as usual.
- Replace 2-node with 3-node.

insert K



search ends here



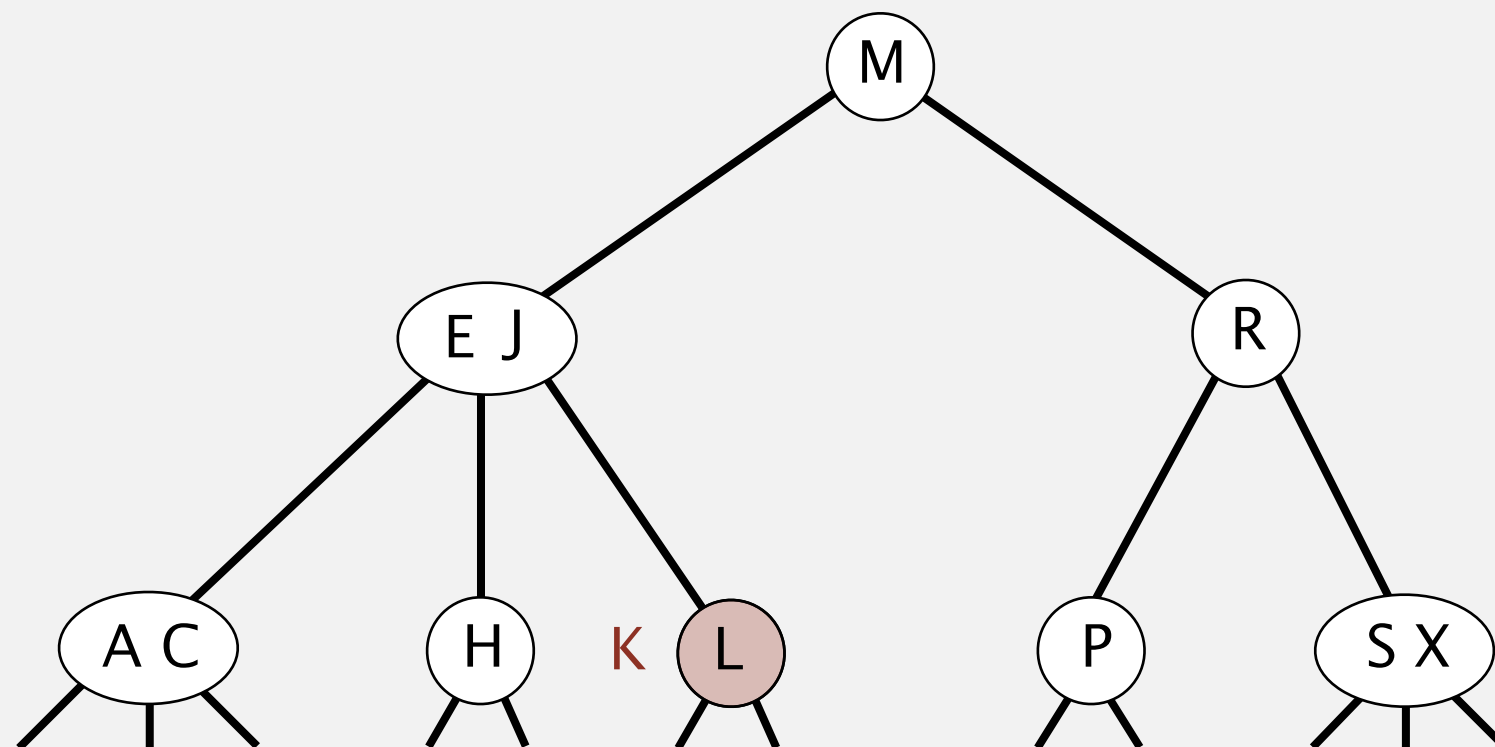
## 2-3 tree demo: insertion

---

Insert into a 2-node at bottom.

- Search for key, as usual.
- Replace 2-node with 3-node.

insert K



replace 2-node with  
3-node containing K

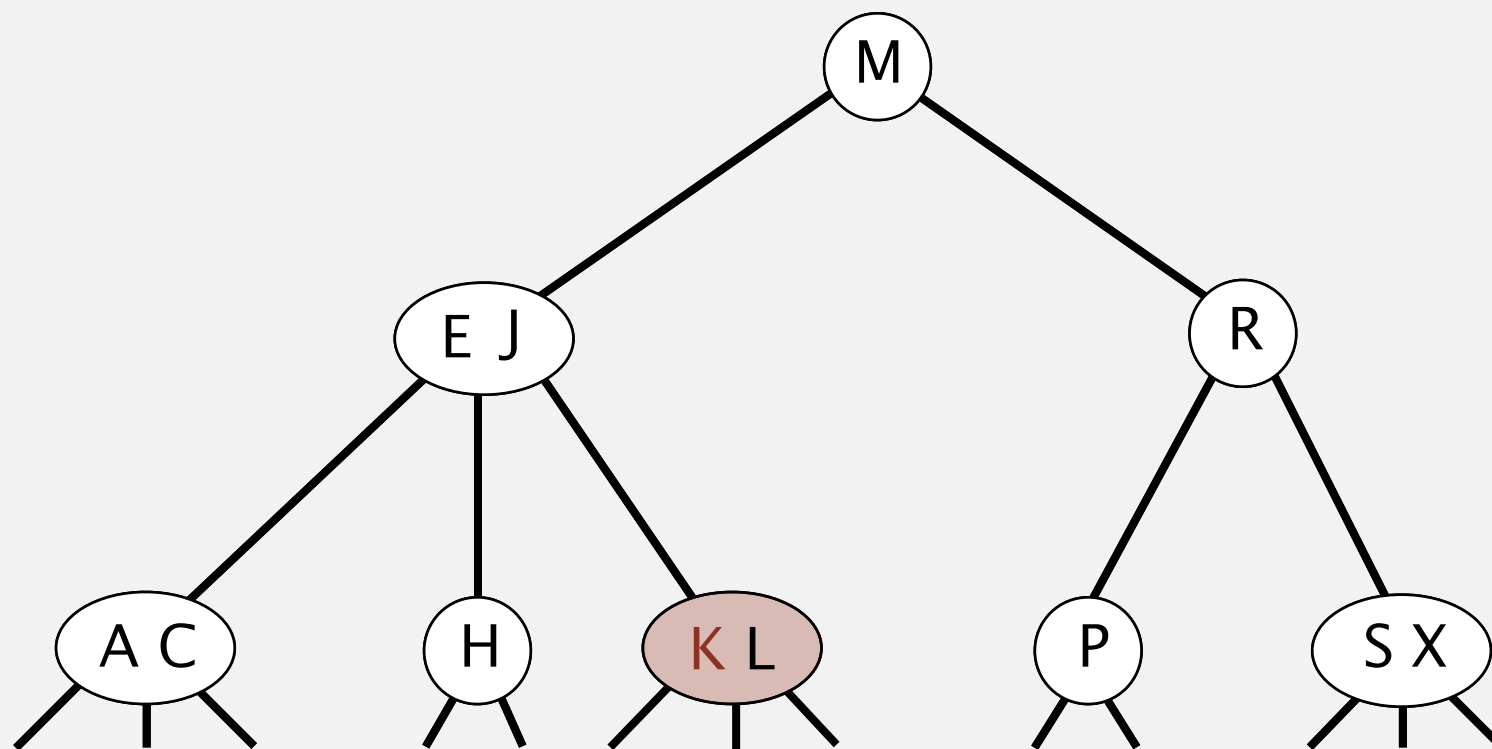
## 2-3 tree demo: insertion

---

Insert into a 2-node at bottom.

- Search for key, as usual.
- Replace 2-node with 3-node.

insert K



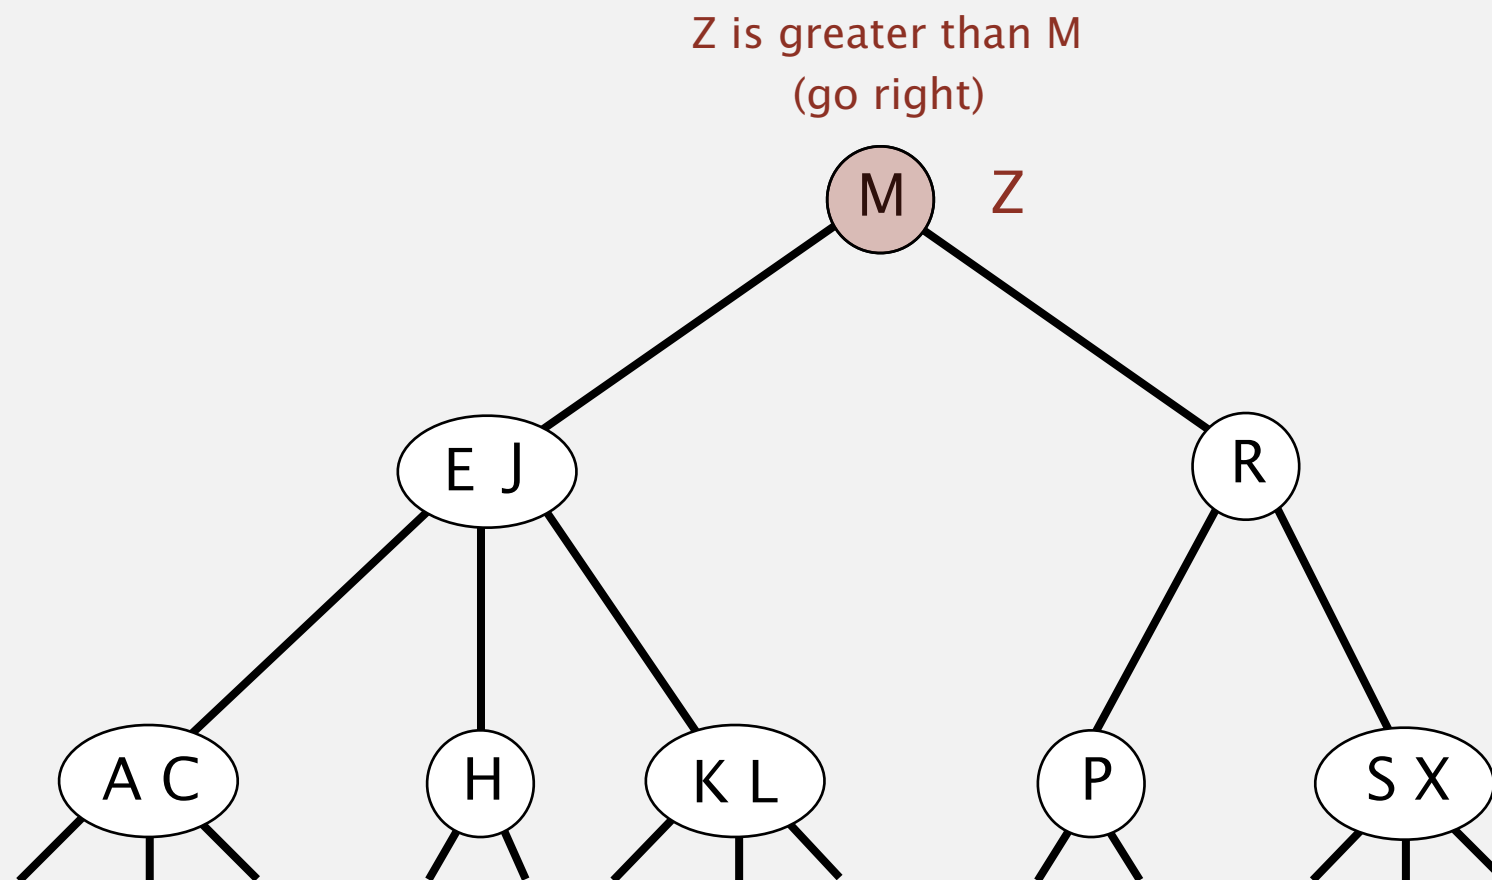
## 2-3 tree demo: insertion

---

Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.

insert Z



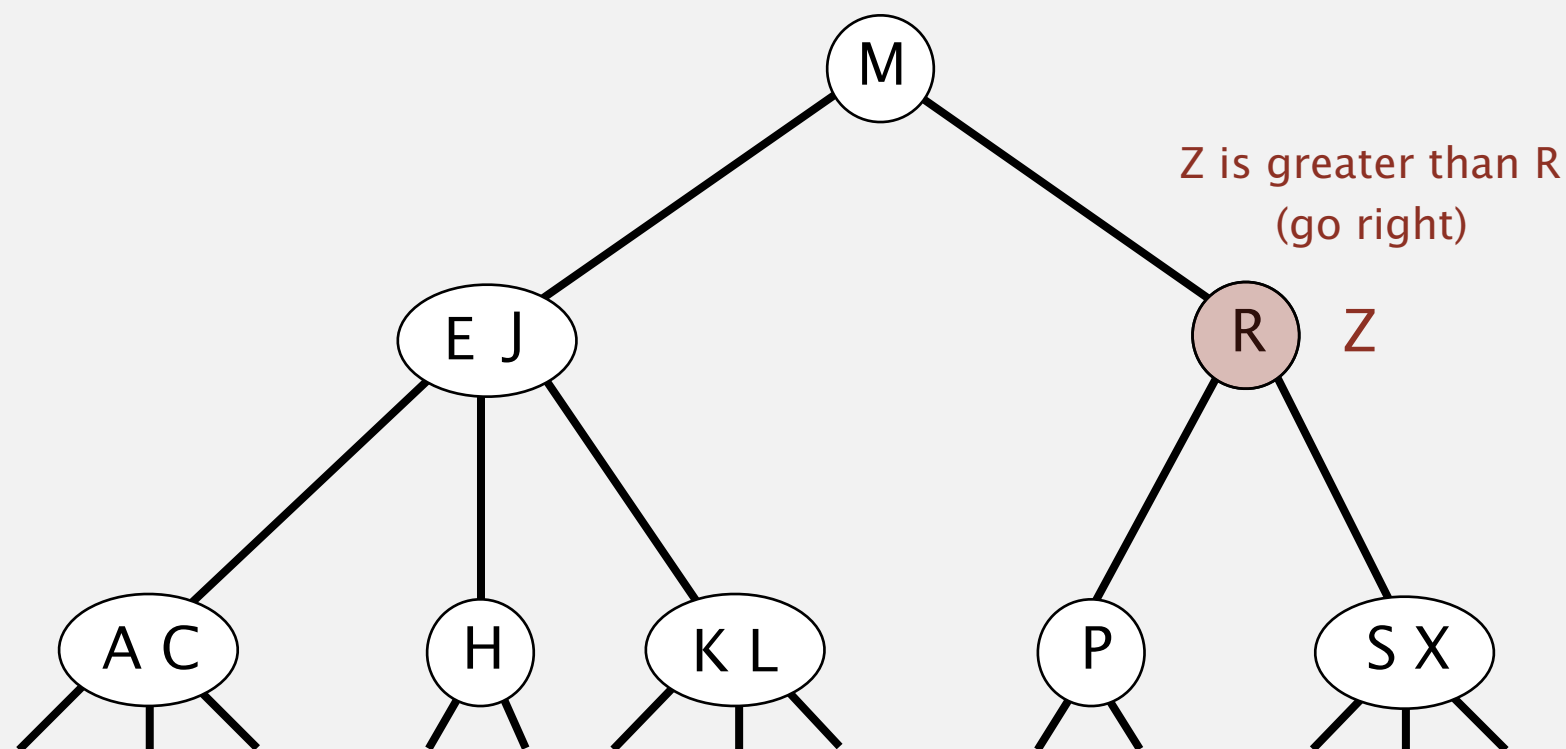
## 2-3 tree demo: insertion

---

Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.

insert Z



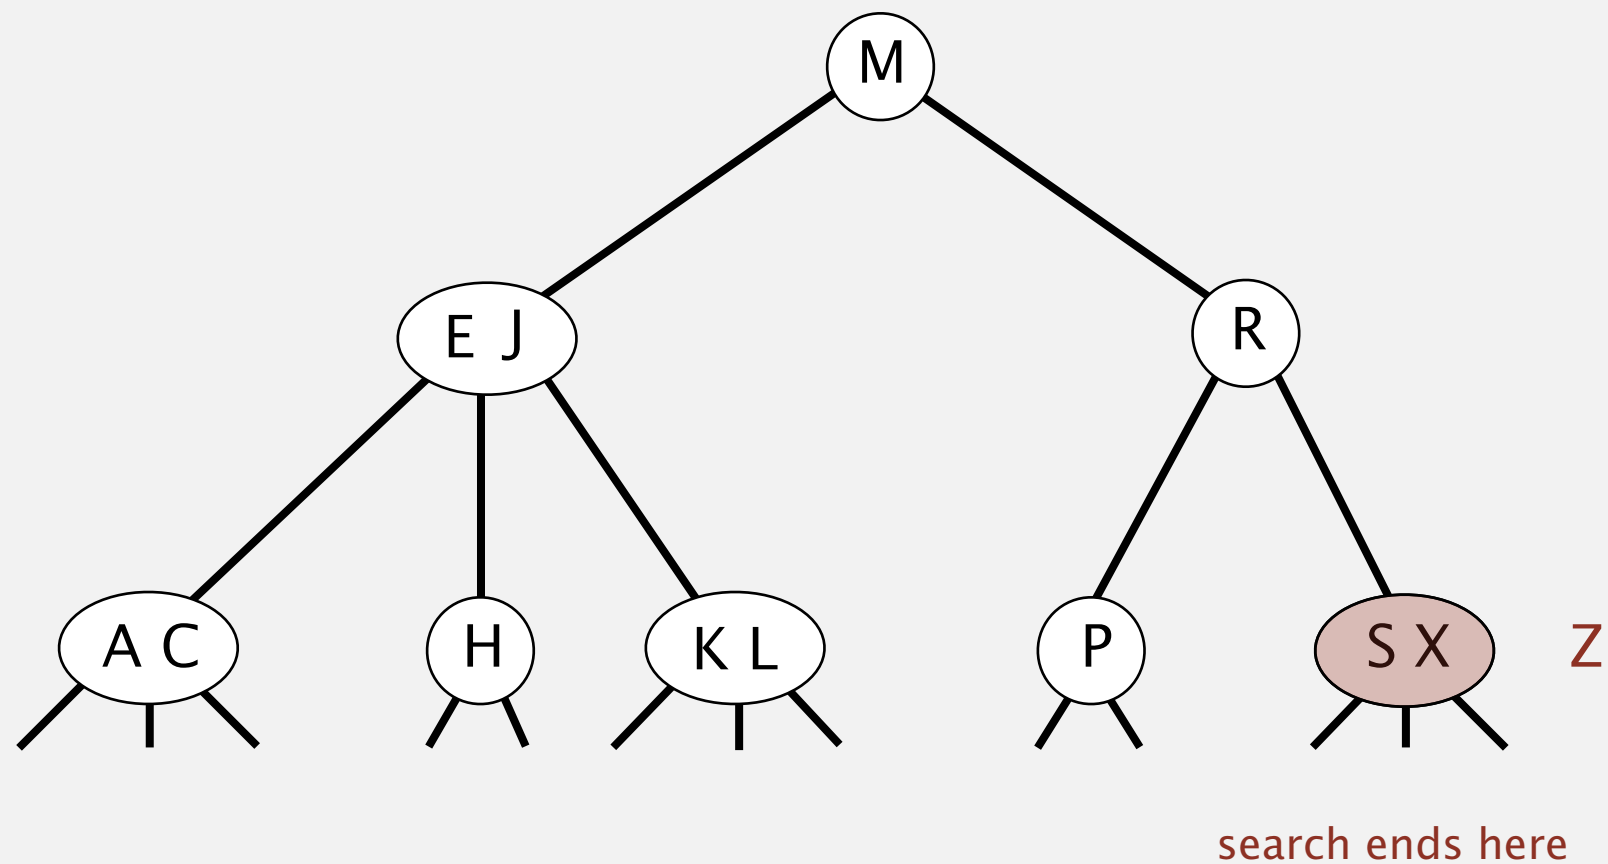
## 2-3 tree demo: insertion

---

Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.

insert Z



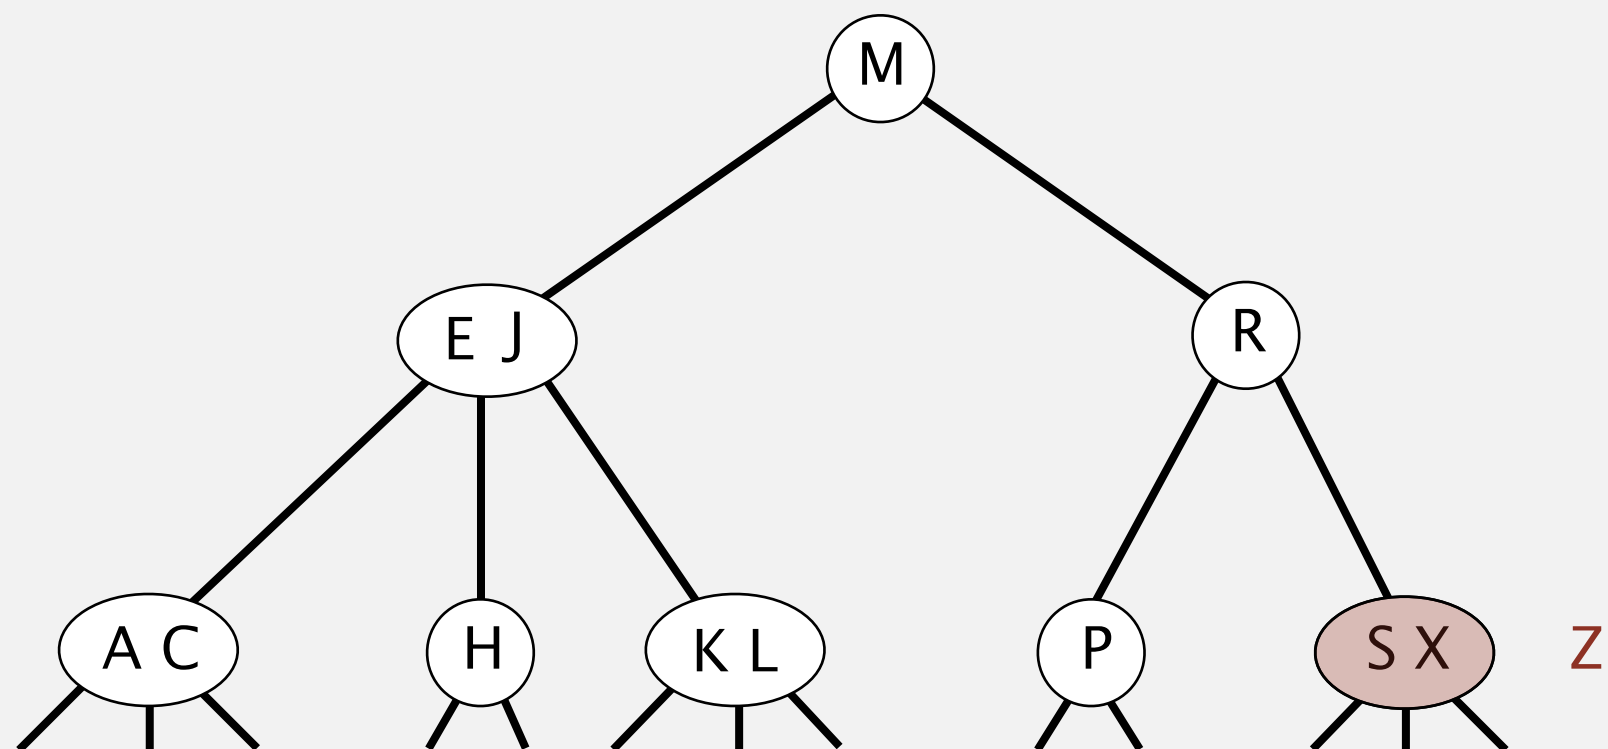
## 2-3 tree demo: insertion

---

Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.

insert Z



replace 3-node with  
temporary 4-node containing Z

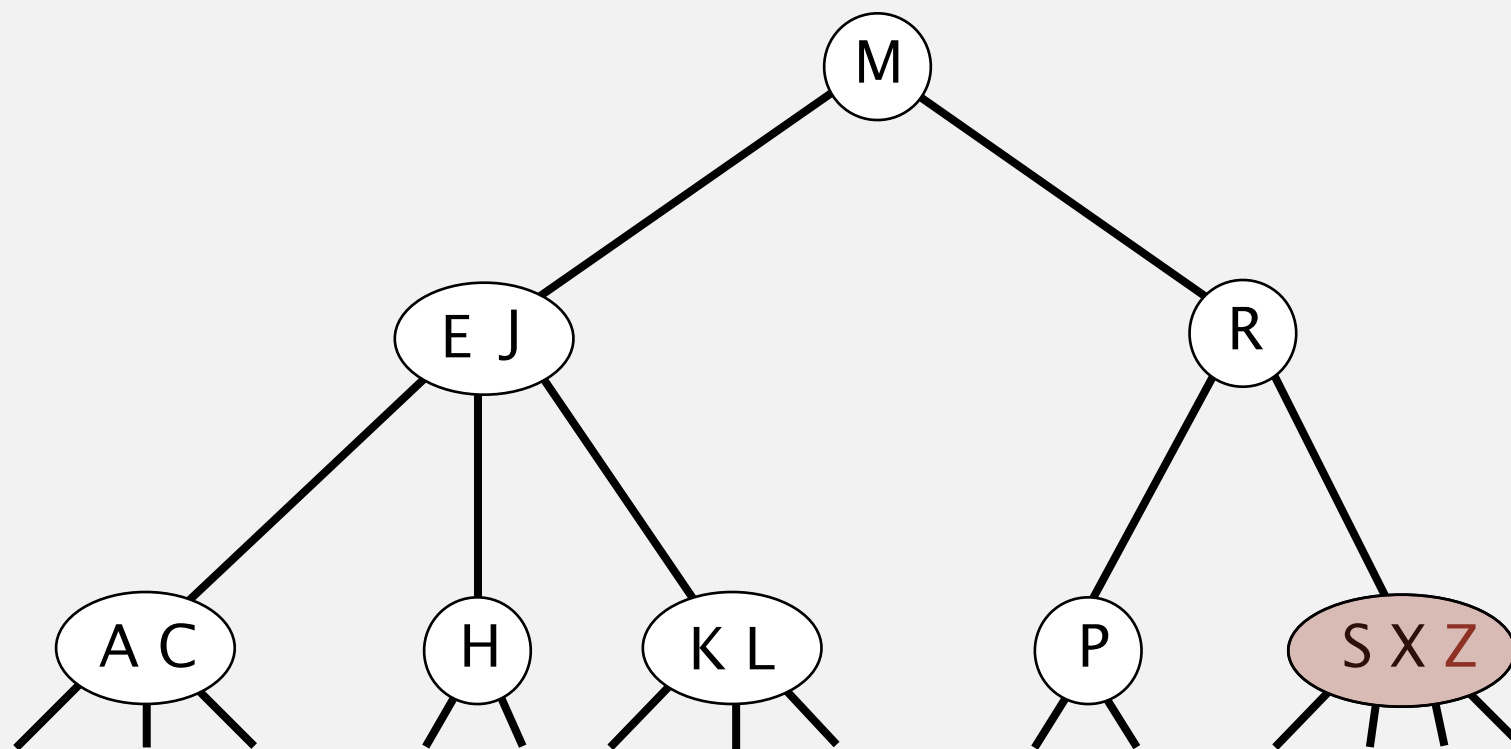
## 2-3 tree demo: insertion

---

Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.

insert Z



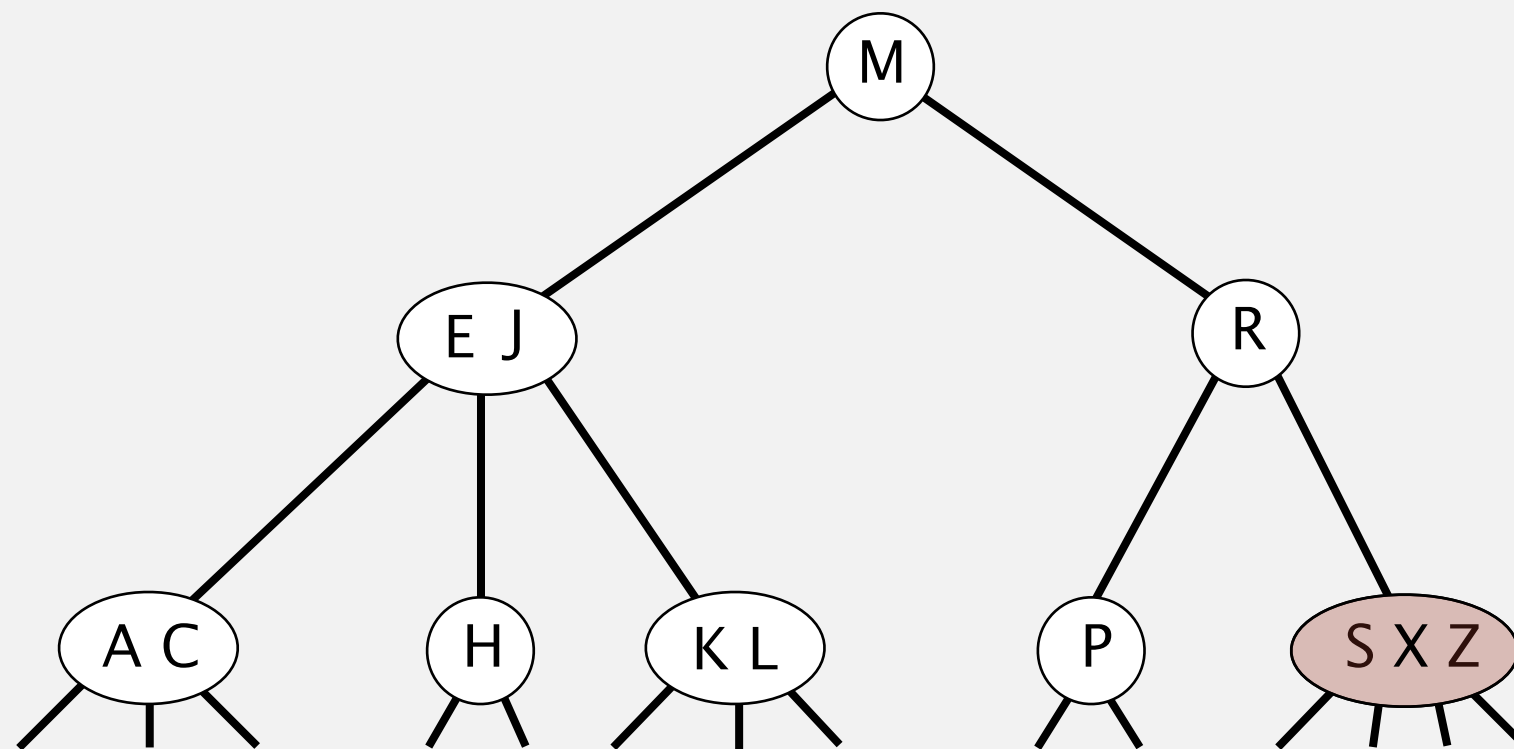
## 2-3 tree demo: insertion

---

Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.

insert Z



split 4-node into two 2-nodes  
(pass middle key to parent)



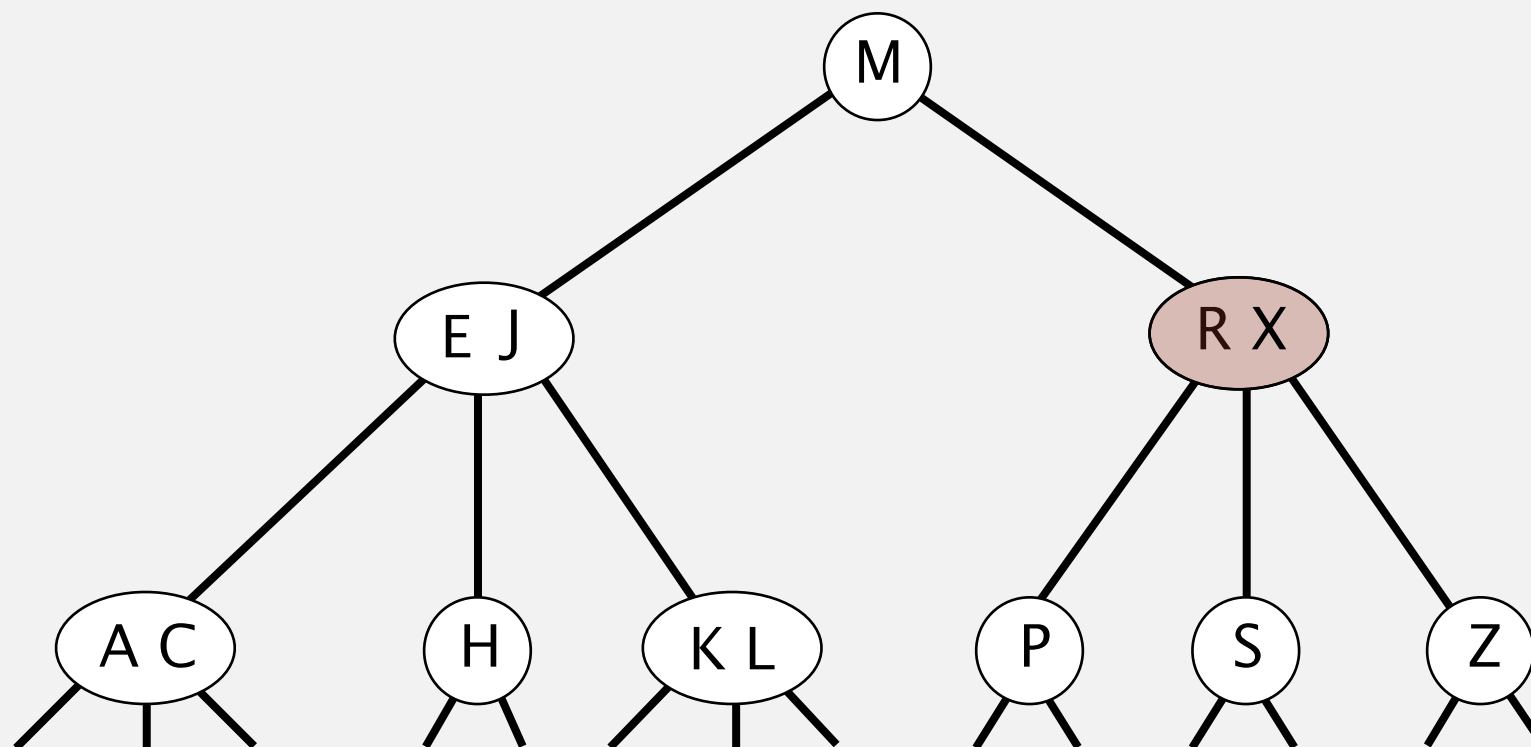
## 2-3 tree demo: insertion

---

Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.

insert Z



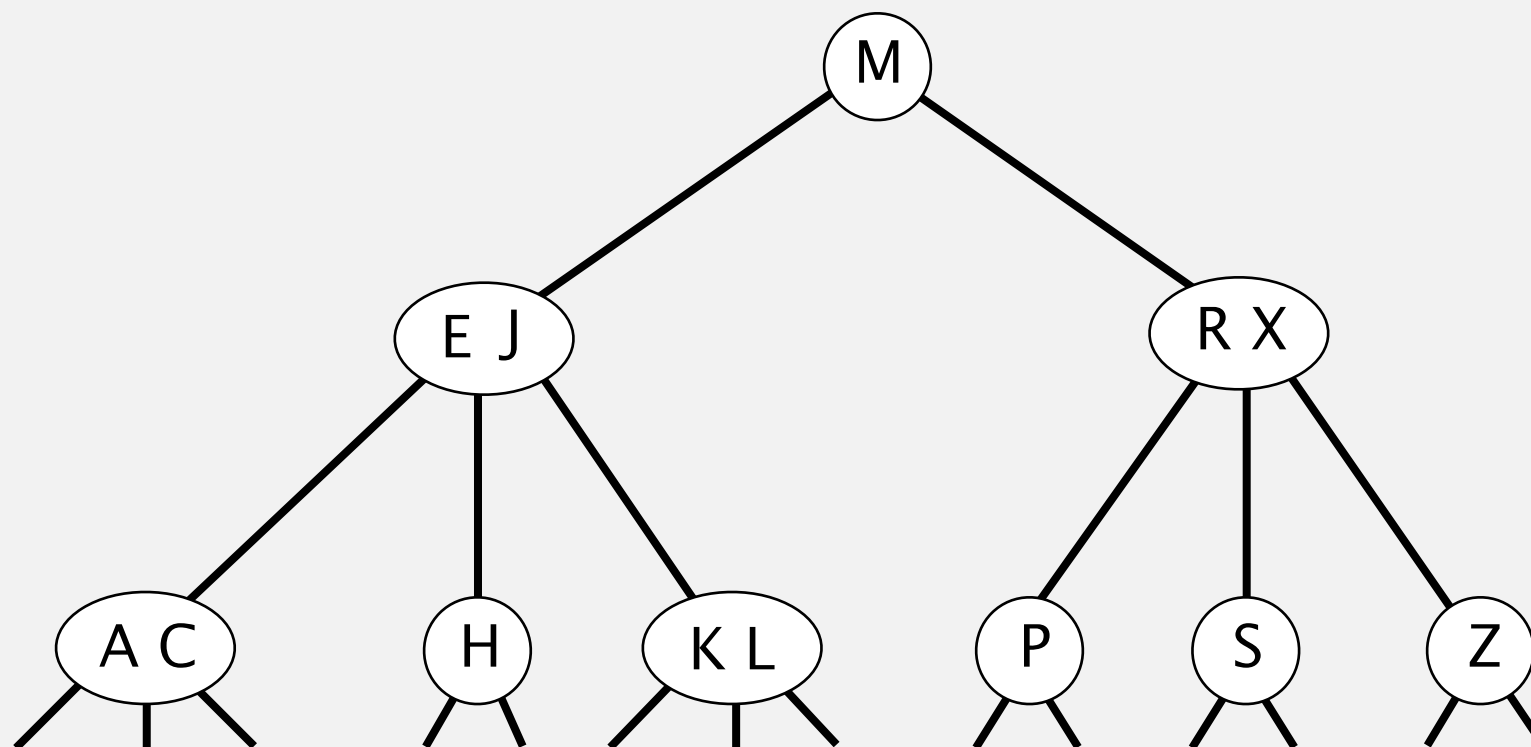
## 2-3 tree demo: insertion

---

Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.

insert Z



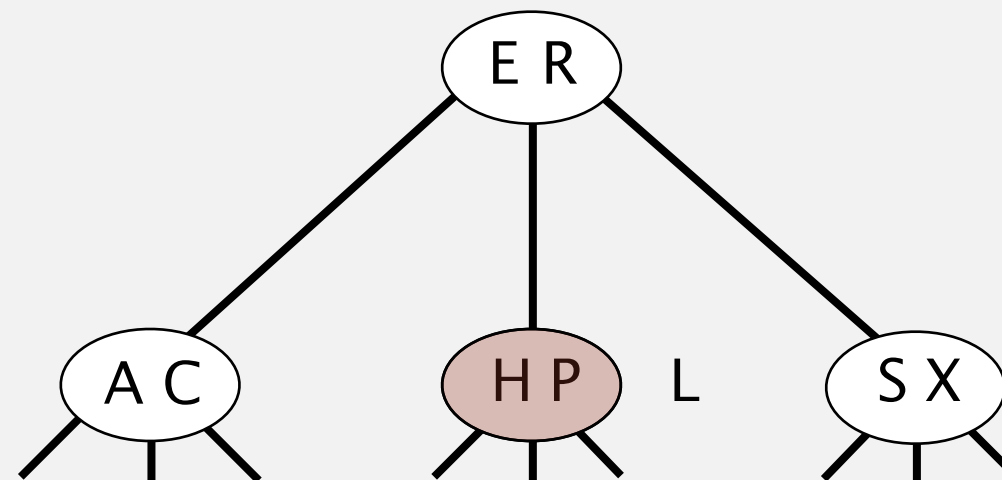
## 2-3 tree demo: insertion

---

### Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.
- Repeat up the tree, as necessary.
- If you reach the root and it's a 4-node, split it into three 2-nodes.

insert L



convert 3-node into 4-node

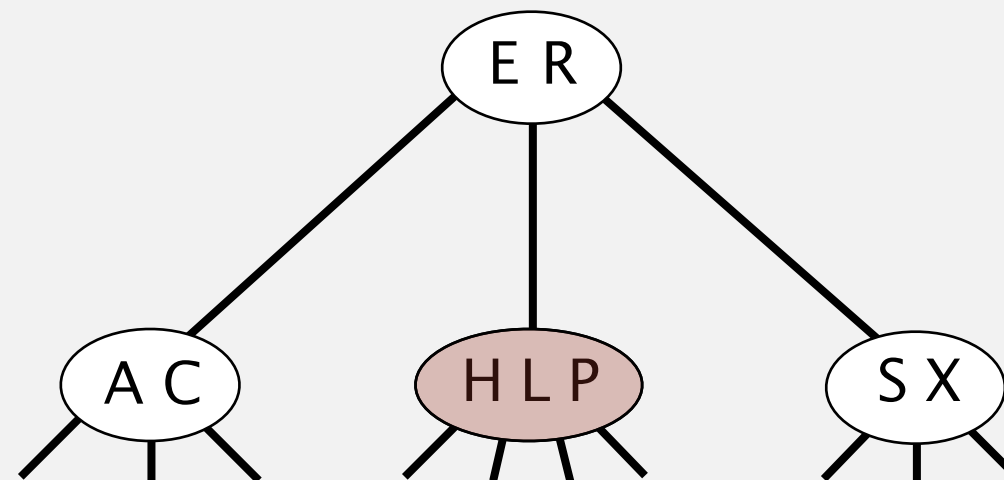
## 2-3 tree demo: insertion

---

### Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.
- Repeat up the tree, as necessary.
- If you reach the root and it's a 4-node, split it into three 2-nodes.

insert L



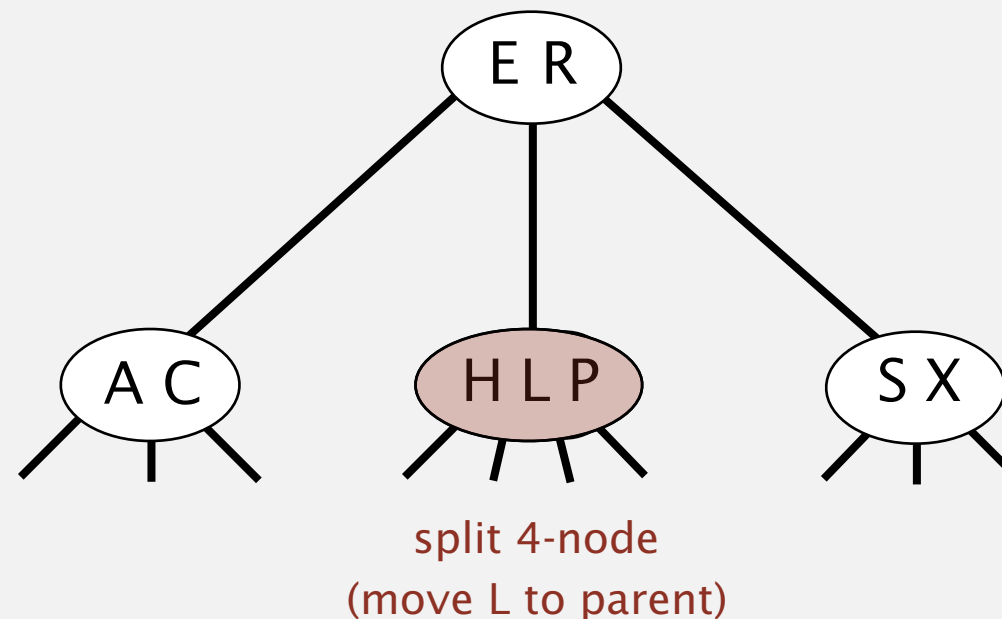
## 2-3 tree demo: insertion

---

### Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.
- Repeat up the tree, as necessary.
- If you reach the root and it's a 4-node, split it into three 2-nodes.

insert L



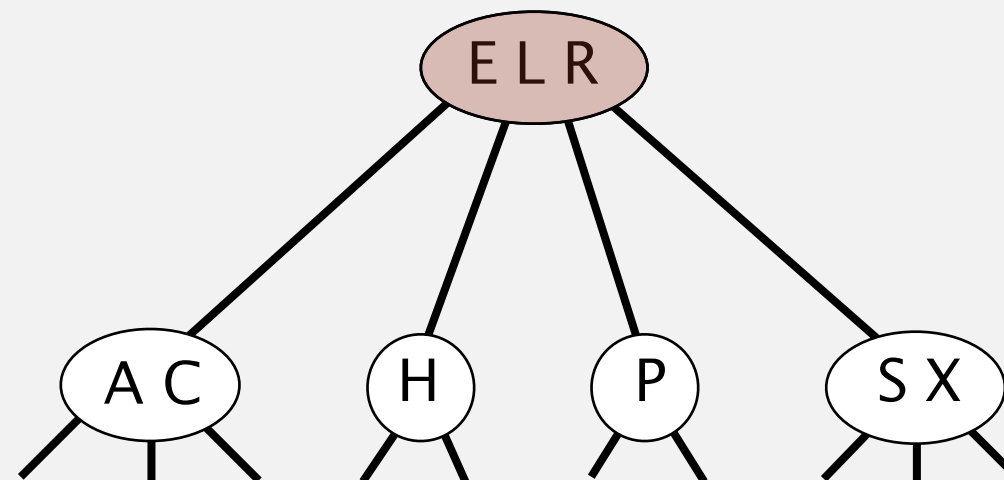
## 2-3 tree demo: insertion

---

### Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.
- Repeat up the tree, as necessary.
- If you reach the root and it's a 4-node, split it into three 2-nodes.

insert L



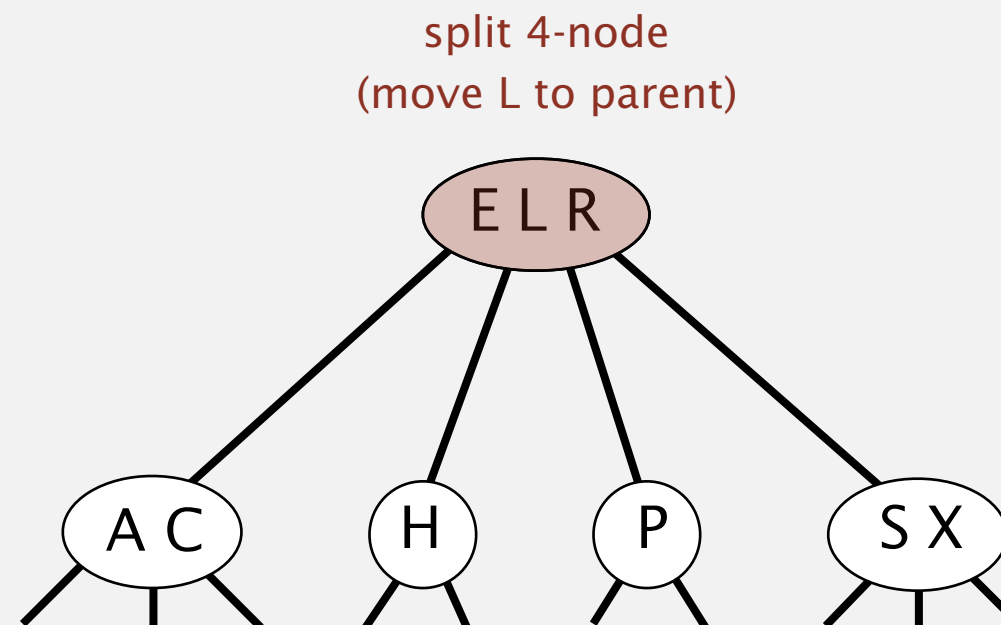
## 2-3 tree demo: insertion

---

### Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.
- Repeat up the tree, as necessary.
- If you reach the root and it's a 4-node, split it into three 2-nodes.

insert L



## 2-3 tree demo: insertion

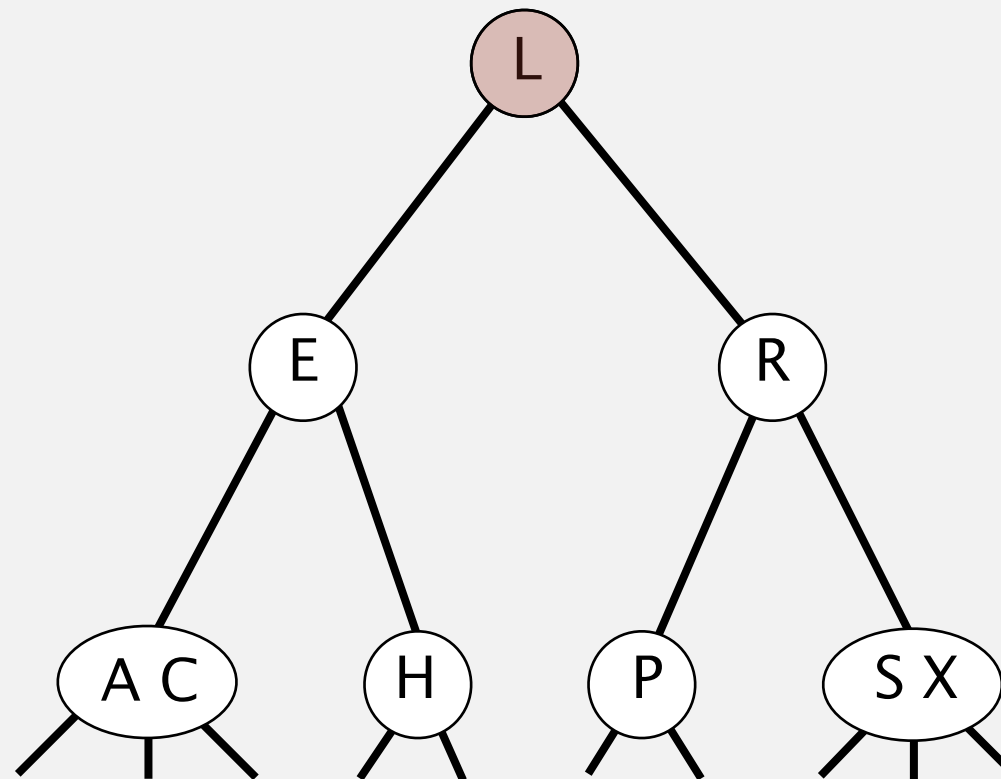
---

### Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.
- Repeat up the tree, as necessary.
- If you reach the root and it's a 4-node, split it into three 2-nodes.

height of tree increases by 1

insert L





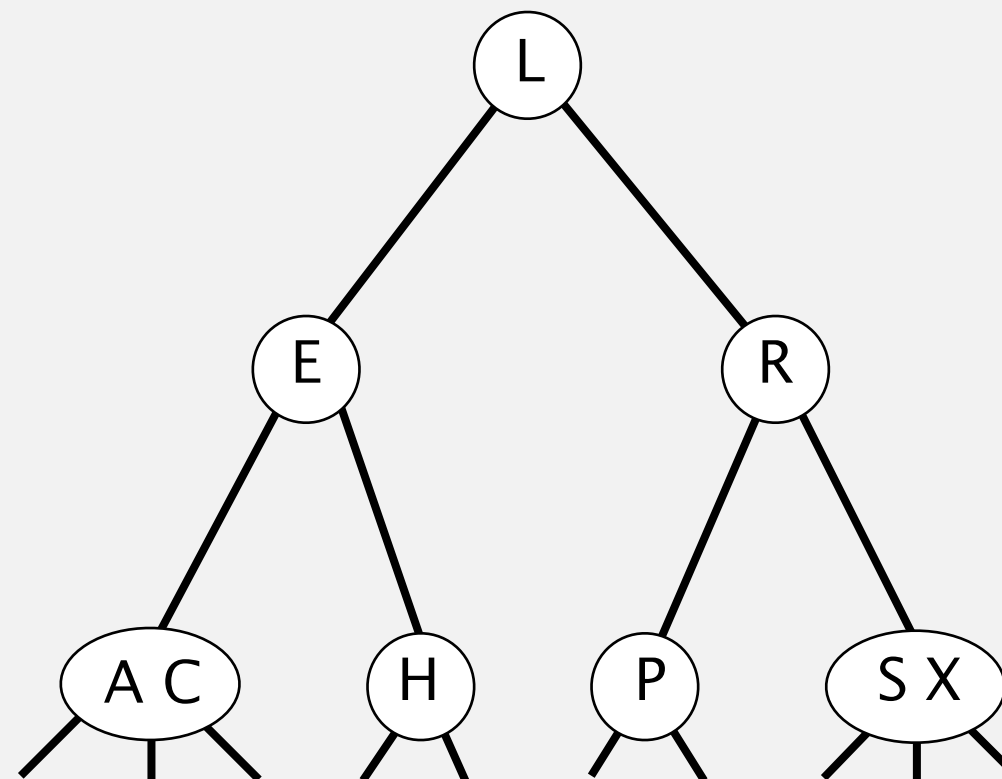
## 2-3 tree demo: insertion

---

Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.
- Repeat up the tree, as necessary.
- If you reach the root and it's a 4-node, split it into three 2-nodes.

insert L

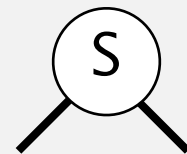


# CONSTRUCTING A 2-3 TREE

## 2-3 tree demo: construction

---

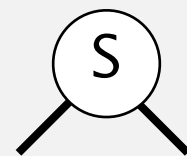
insert S



# 2-3 tree demo: construction

---

2-3 tree



# 2-3 tree demo: construction

---

insert E

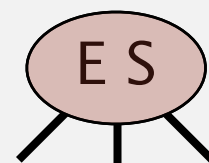


convert 2-node into 3-node

## 2-3 tree demo: construction

---

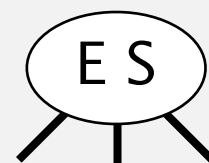
insert E



# 2-3 tree demo: construction

---

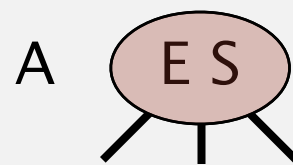
2-3 tree



## 2-3 tree demo: construction

---

insert A



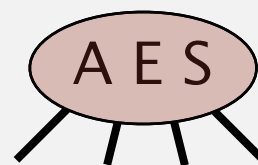
convert 3-node into 4-node



## 2-3 tree demo: construction

---

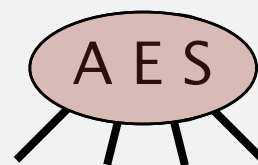
insert A



## 2-3 tree demo: construction

---

insert A

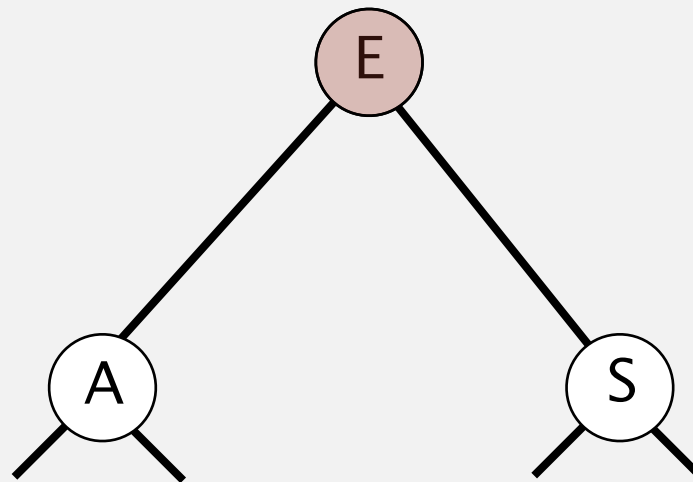


split 4-node  
(move E to parent)

## 2-3 tree demo: construction

---

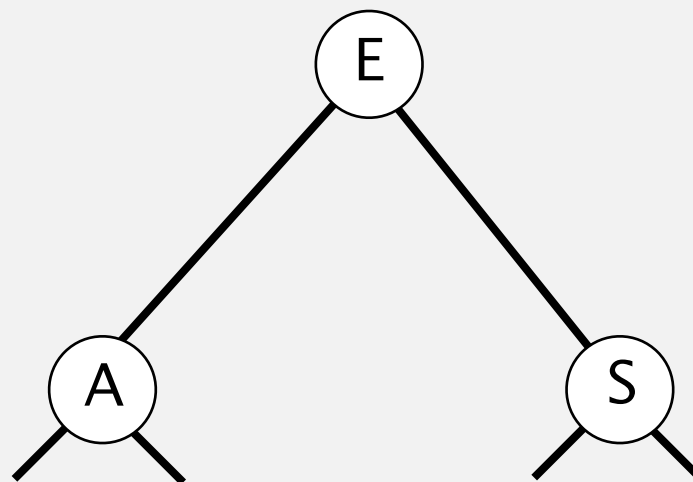
insert A



# 2-3 tree demo: construction

---

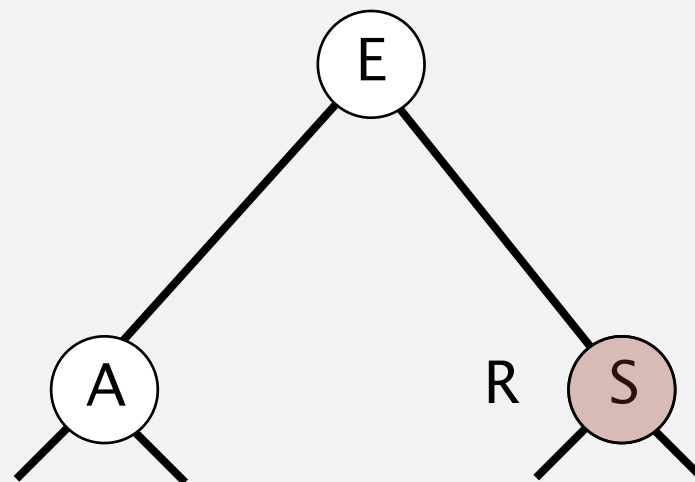
2-3 tree



## 2-3 tree demo: construction

---

insert R

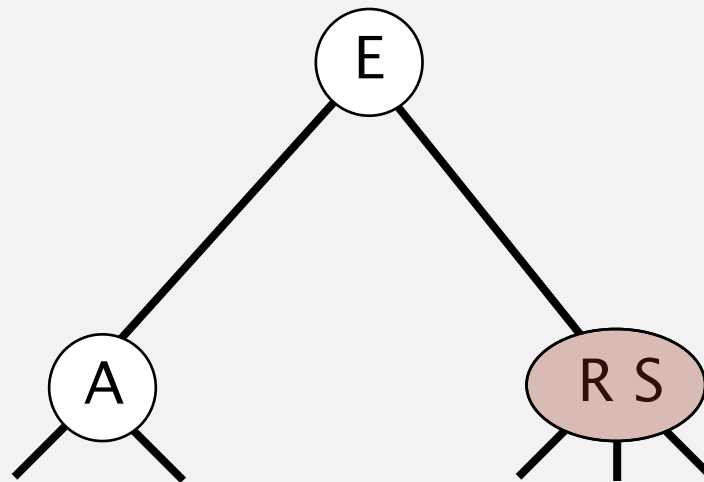


convert 2-node into 3-node

## 2-3 tree demo: construction

---

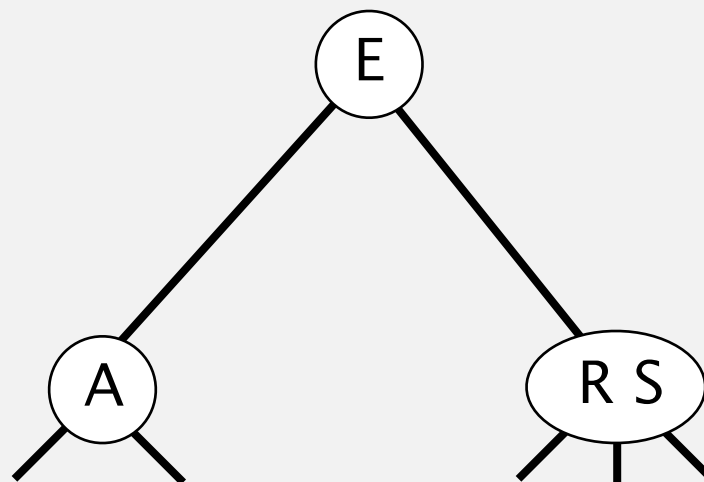
insert R



## 2-3 tree demo: construction

---

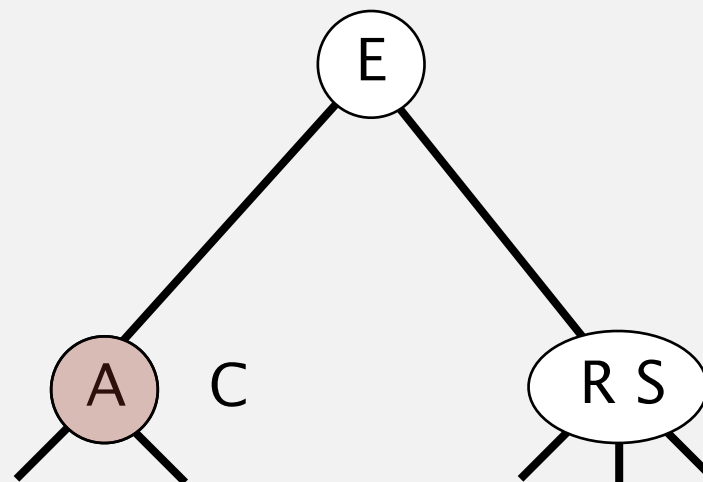
2-3 tree



## 2-3 tree demo: construction

---

insert C



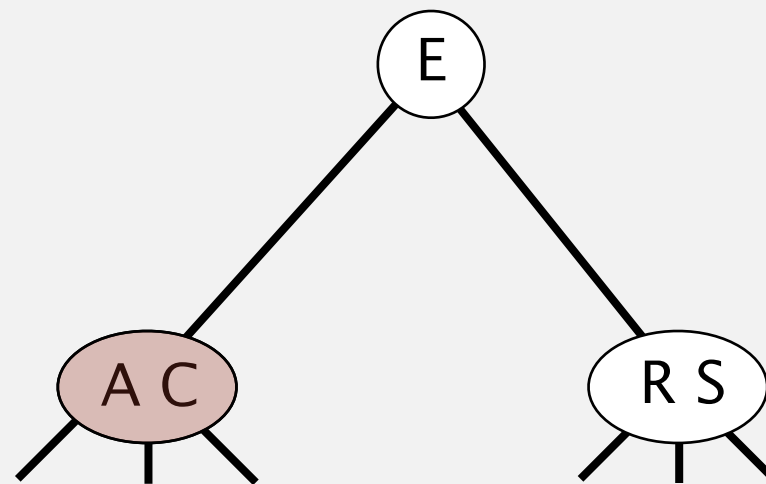
convert 2-node into 3-node



## 2-3 tree demo: construction

---

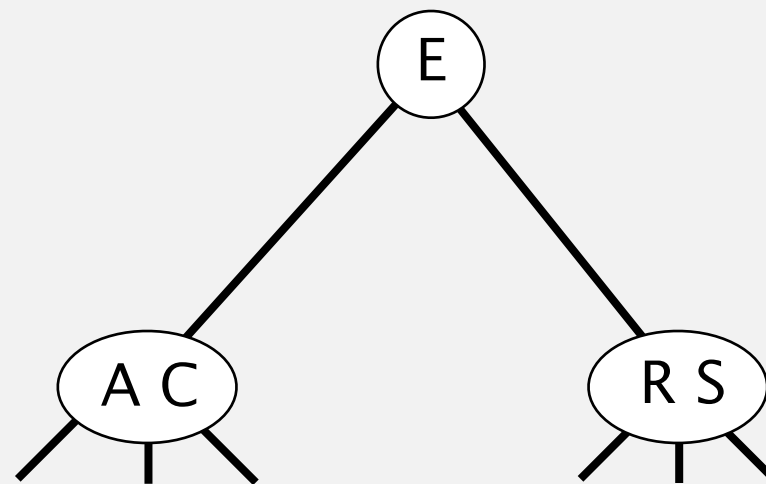
insert C



## 2-3 tree demo: construction

---

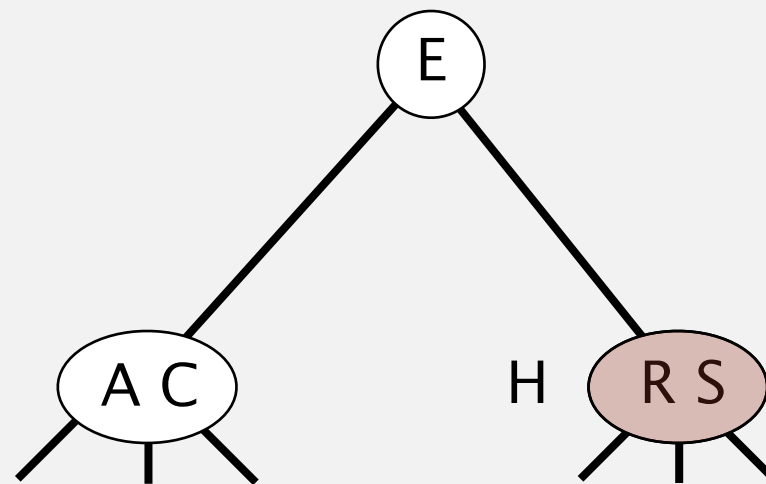
2-3 tree



## 2-3 tree demo: construction

---

insert H

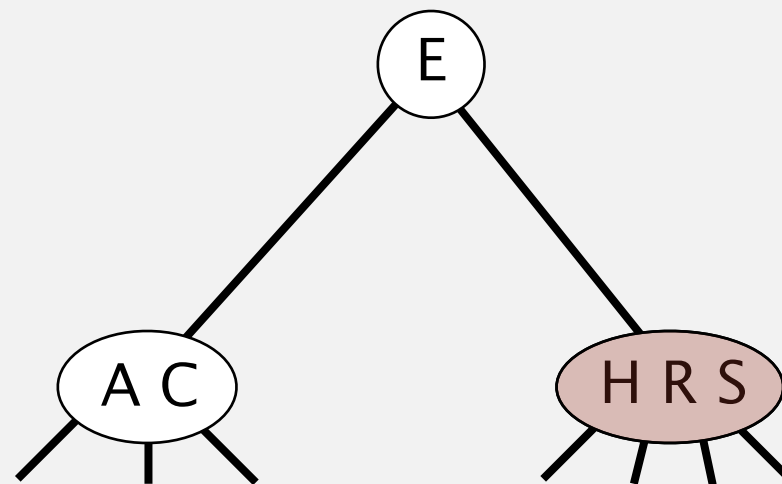


convert 3-node into 4-node

## 2-3 tree demo: construction

---

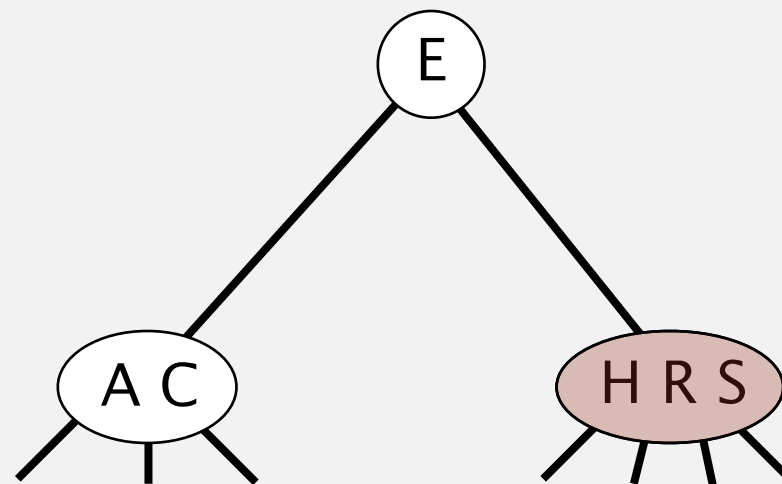
insert H



## 2-3 tree demo: construction

---

insert H

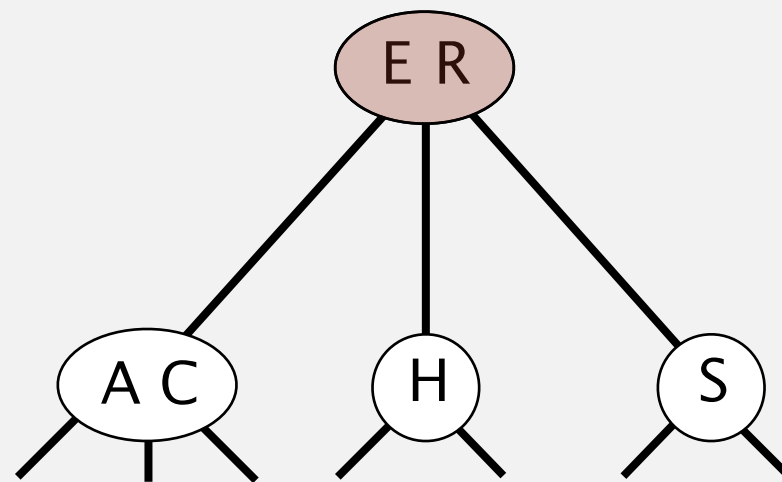


split 4-node  
(move R to parent)

## 2-3 tree demo: construction

---

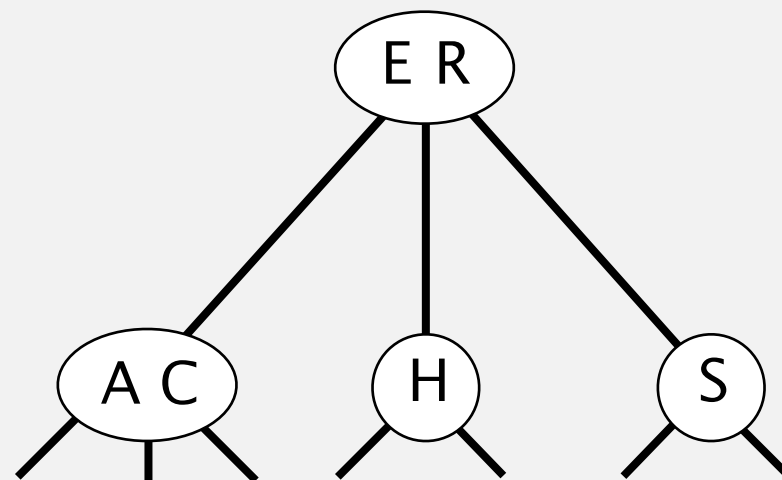
insert H



## 2-3 tree demo: construction

---

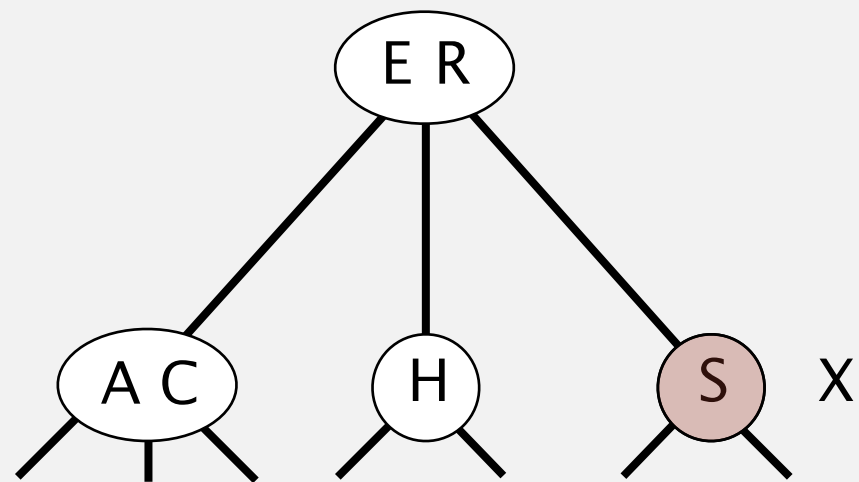
2-3 tree



## 2-3 tree demo: construction

---

insert X



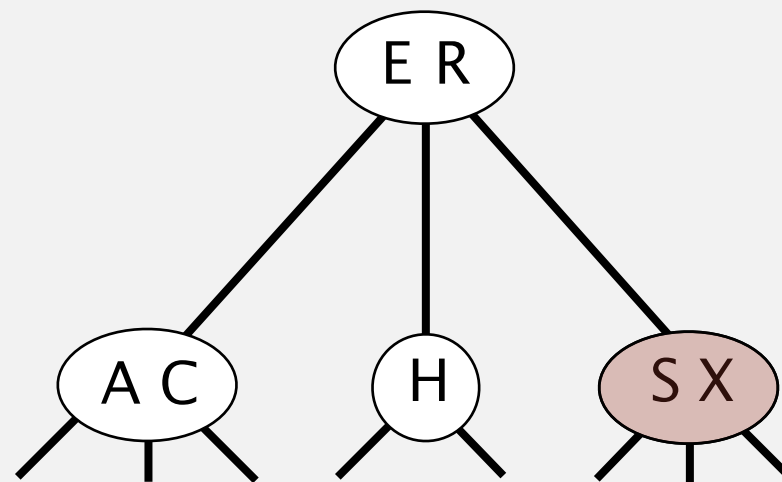
convert 2-node into 3-node



## 2-3 tree demo: construction

---

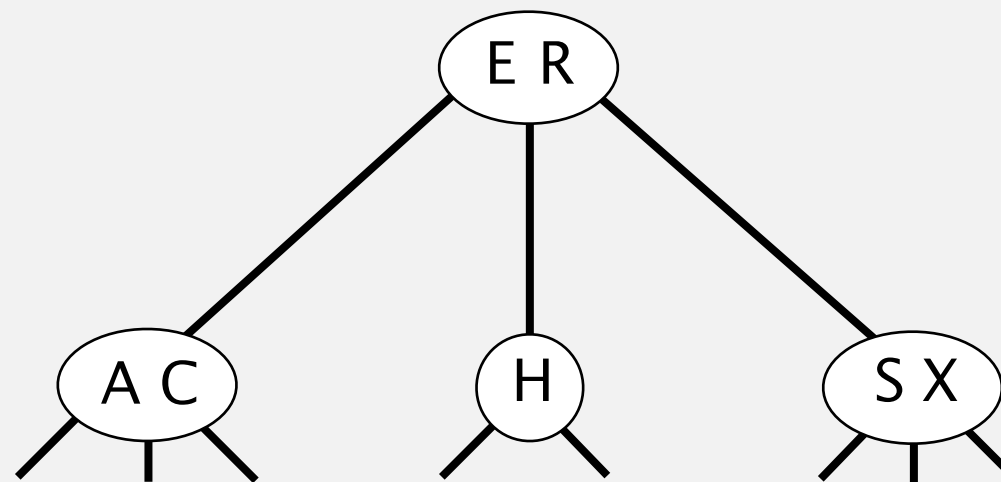
insert X



## 2-3 tree demo: construction

---

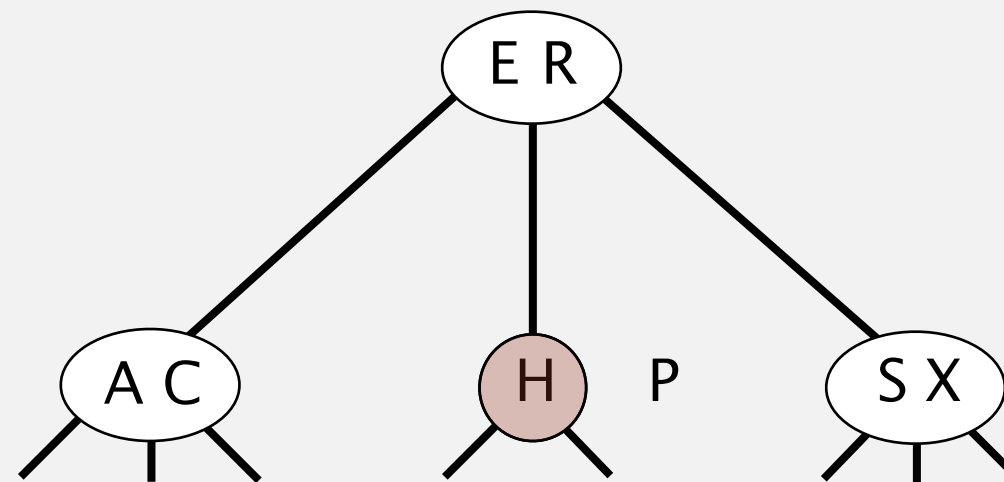
2-3 tree



## 2-3 tree demo: construction

---

insert P

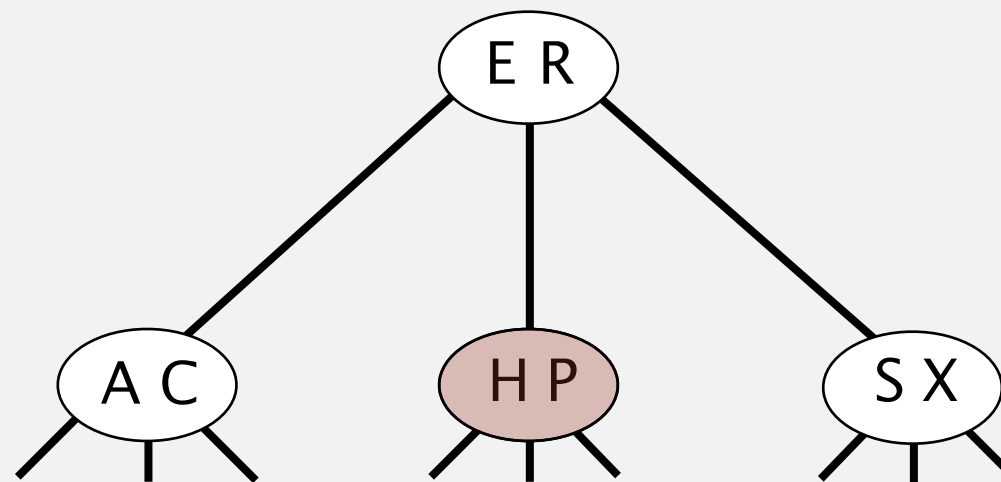


convert 2-node into 3-node

## 2-3 tree demo: construction

---

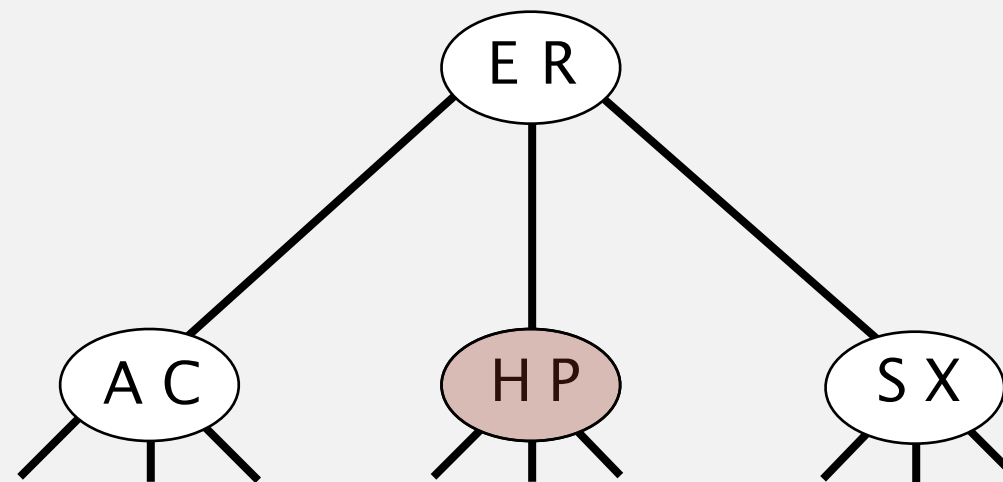
insert P



# Quiz

---

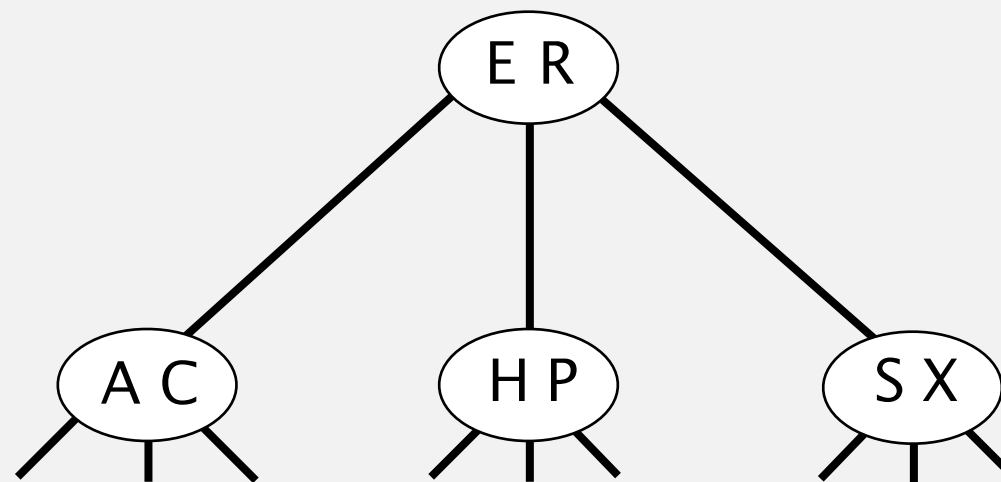
- What is the result of inserting L into the following tree?



## 2-3 tree demo: construction

---

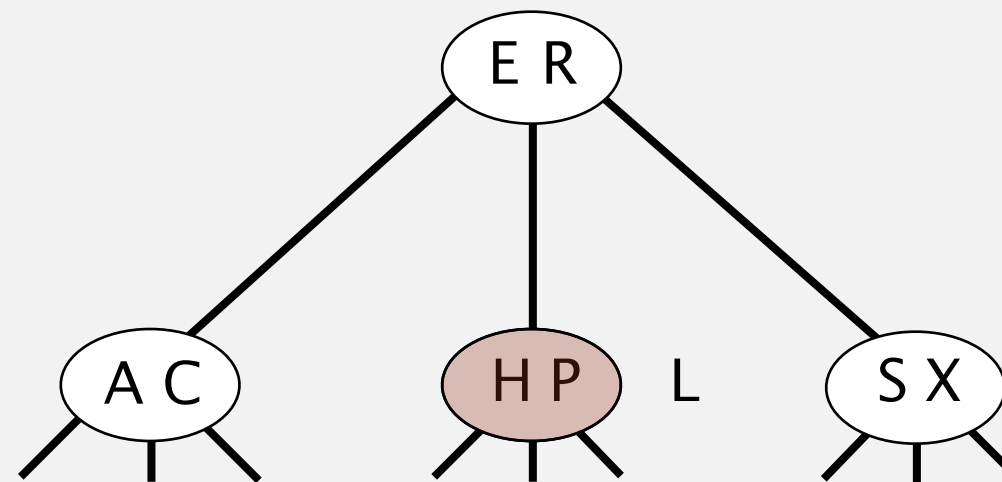
2-3 tree



## 2-3 tree demo: construction

---

insert L

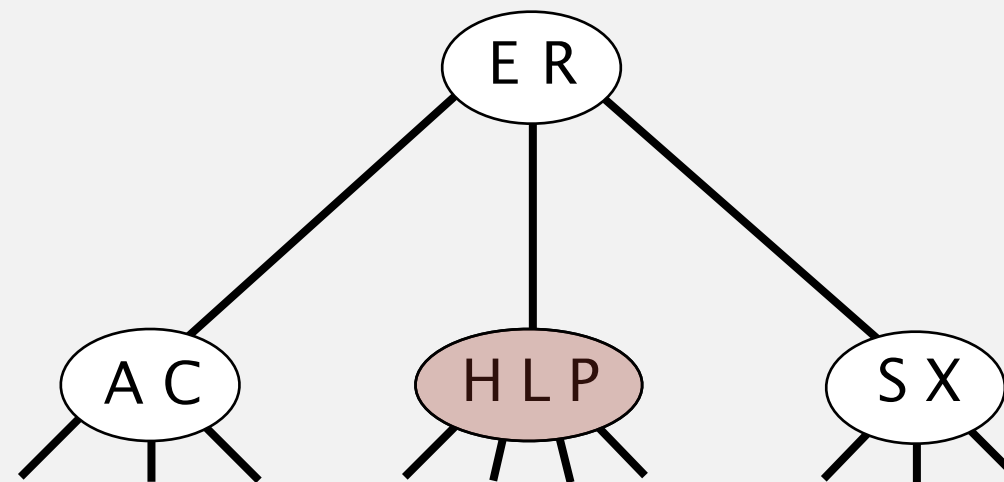


convert 3-node into 4-node

## 2-3 tree demo: construction

---

insert L

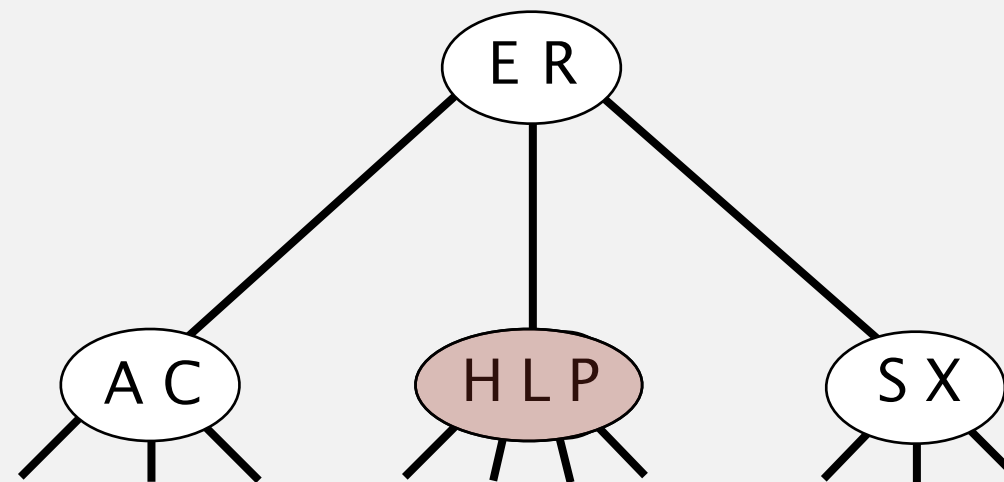




## 2-3 tree demo: construction

---

insert L

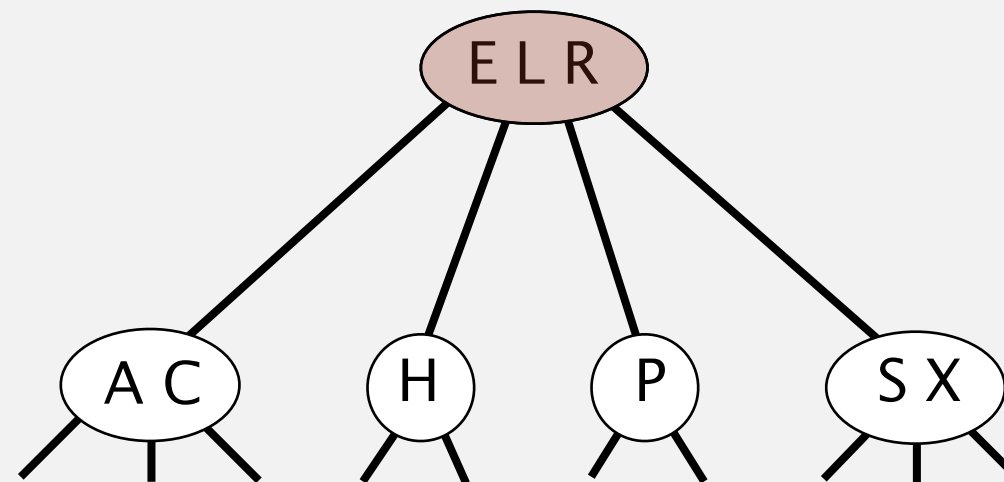


split 4-node  
(move L to parent)

## 2-3 tree demo: construction

---

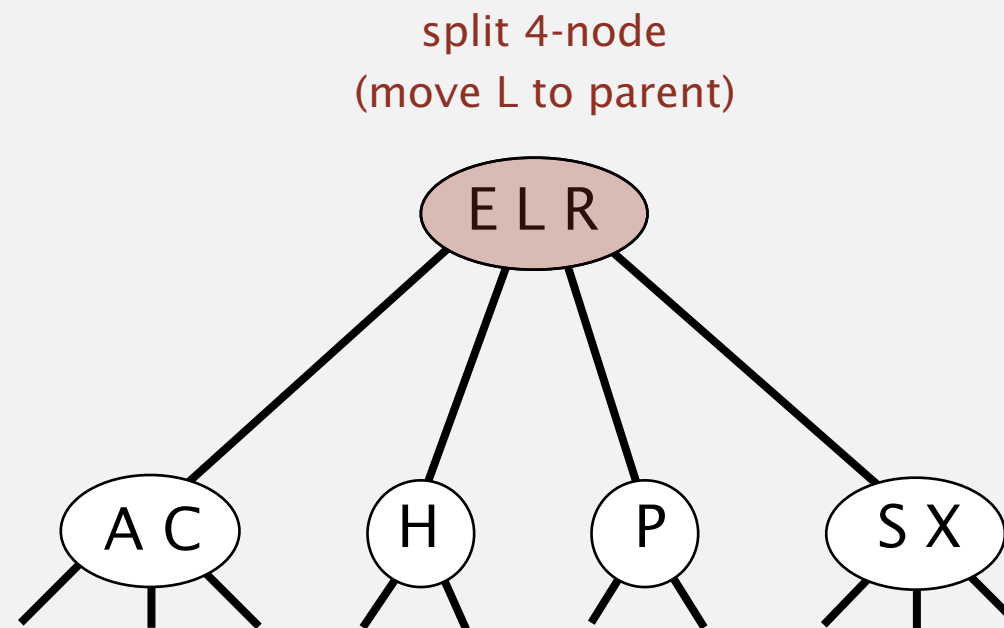
insert L



## 2-3 tree demo: construction

---

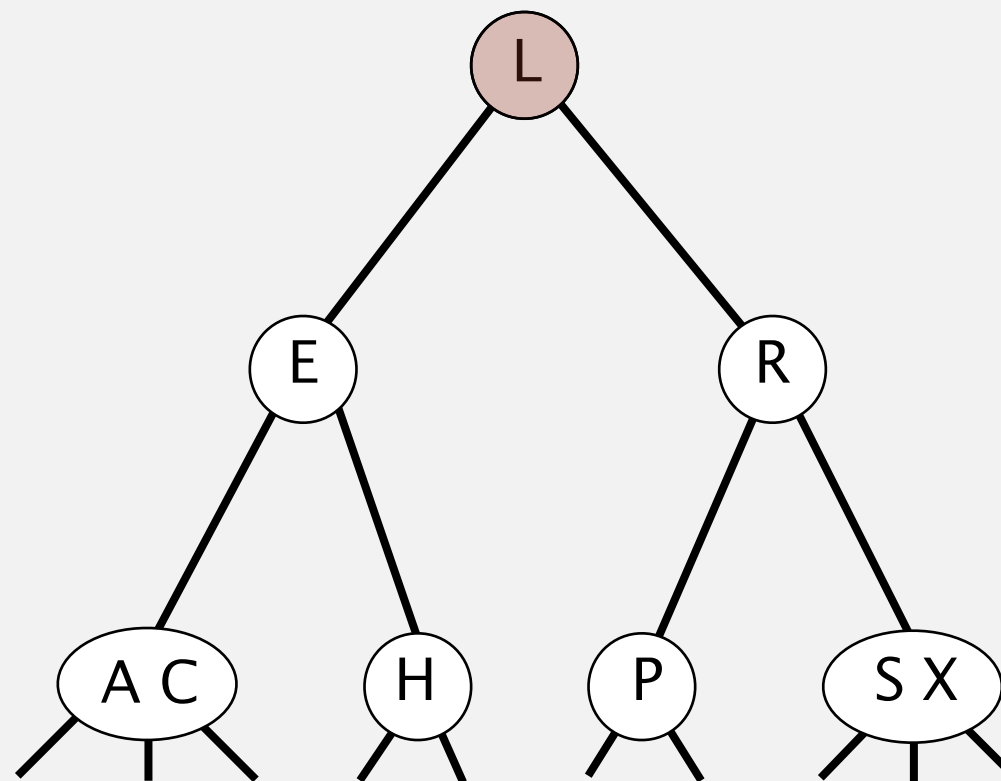
insert L



## 2-3 tree demo: construction

---

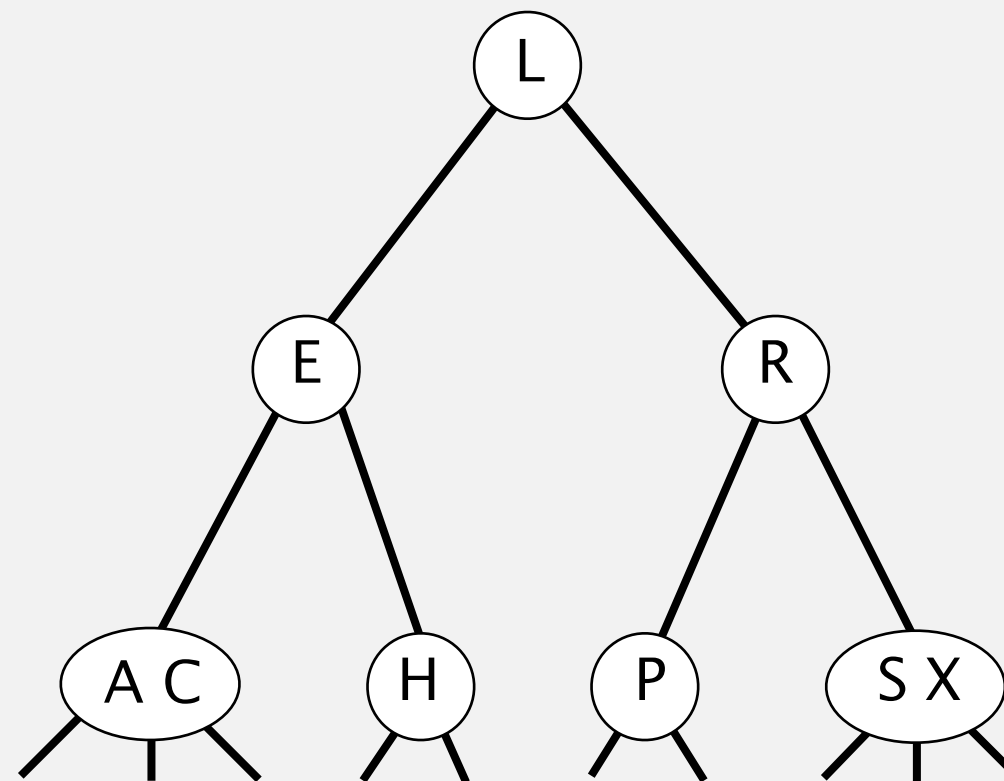
insert L



## 2-3 tree demo: construction

---

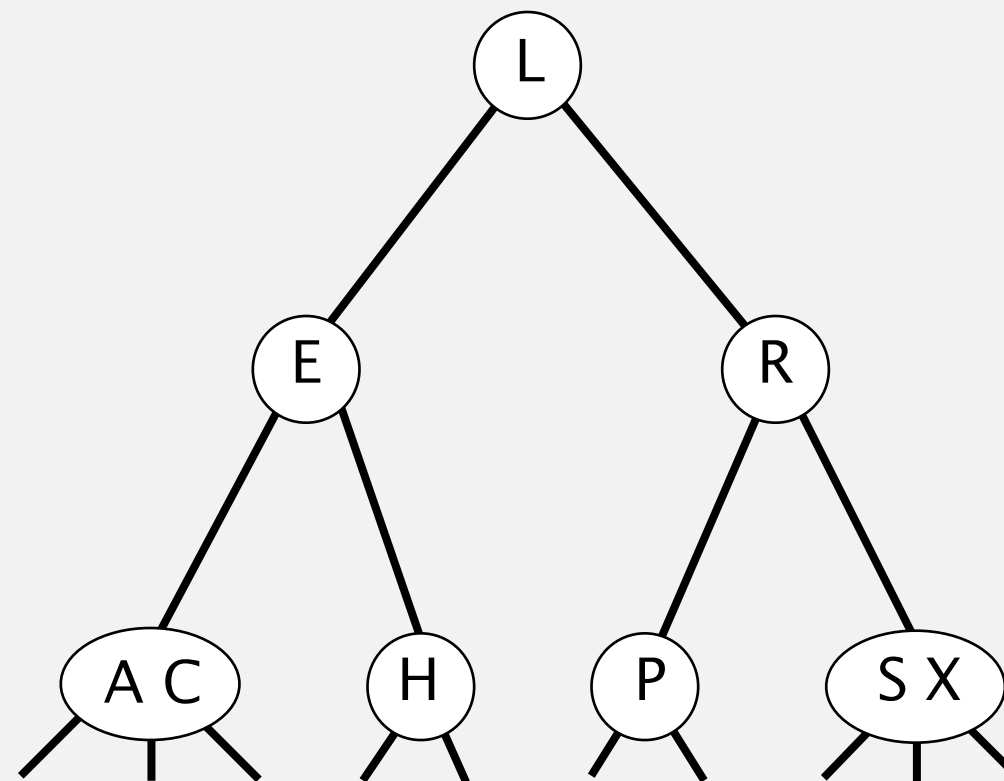
2-3 tree



# Quiz

---

What is the result of inserting G into the following tree



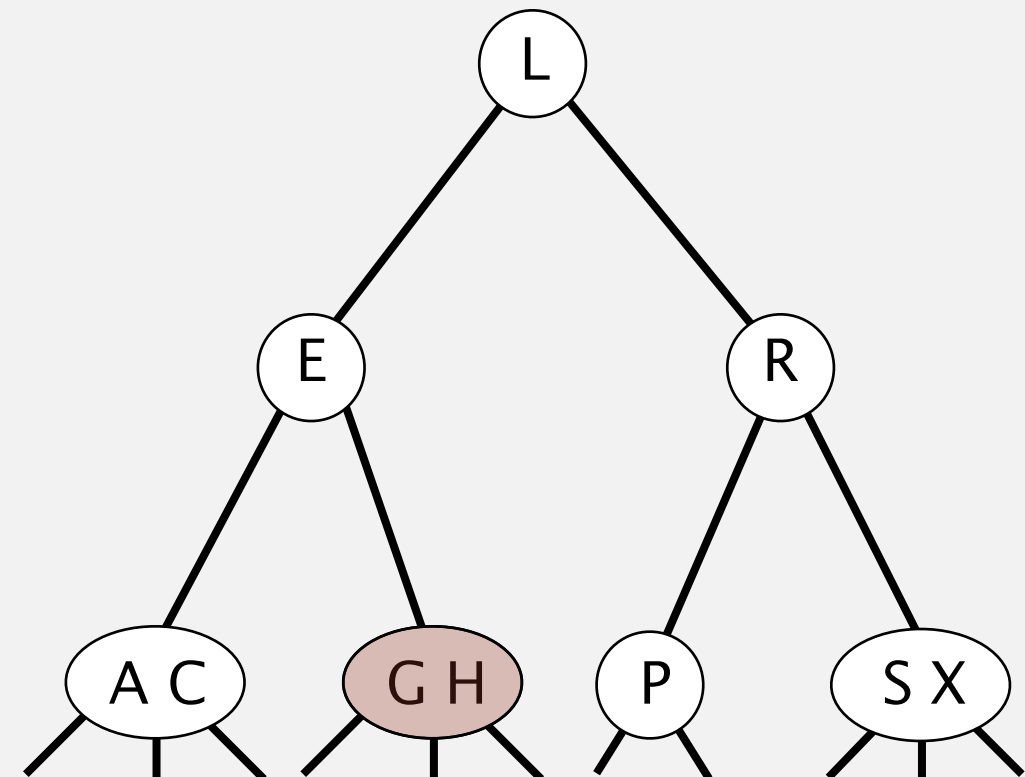
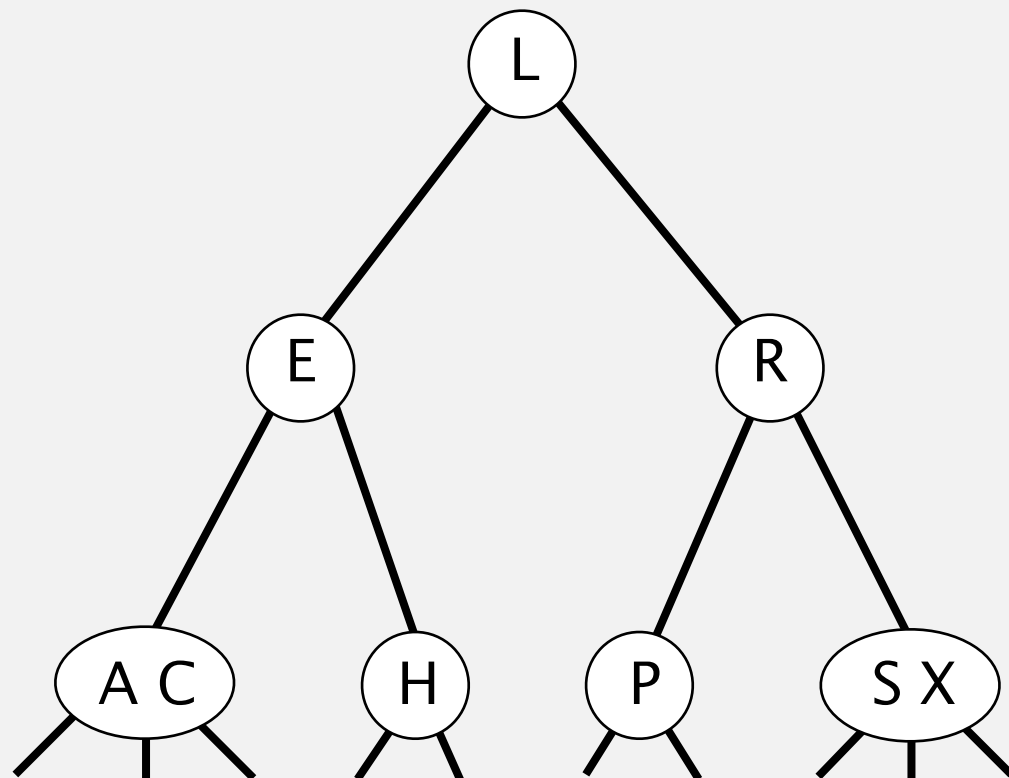
# Insertion into a 2-3 tree

---

## Insertion into a 2-node at bottom.

- Add new key to 2-node to create a 3-node.

insert G

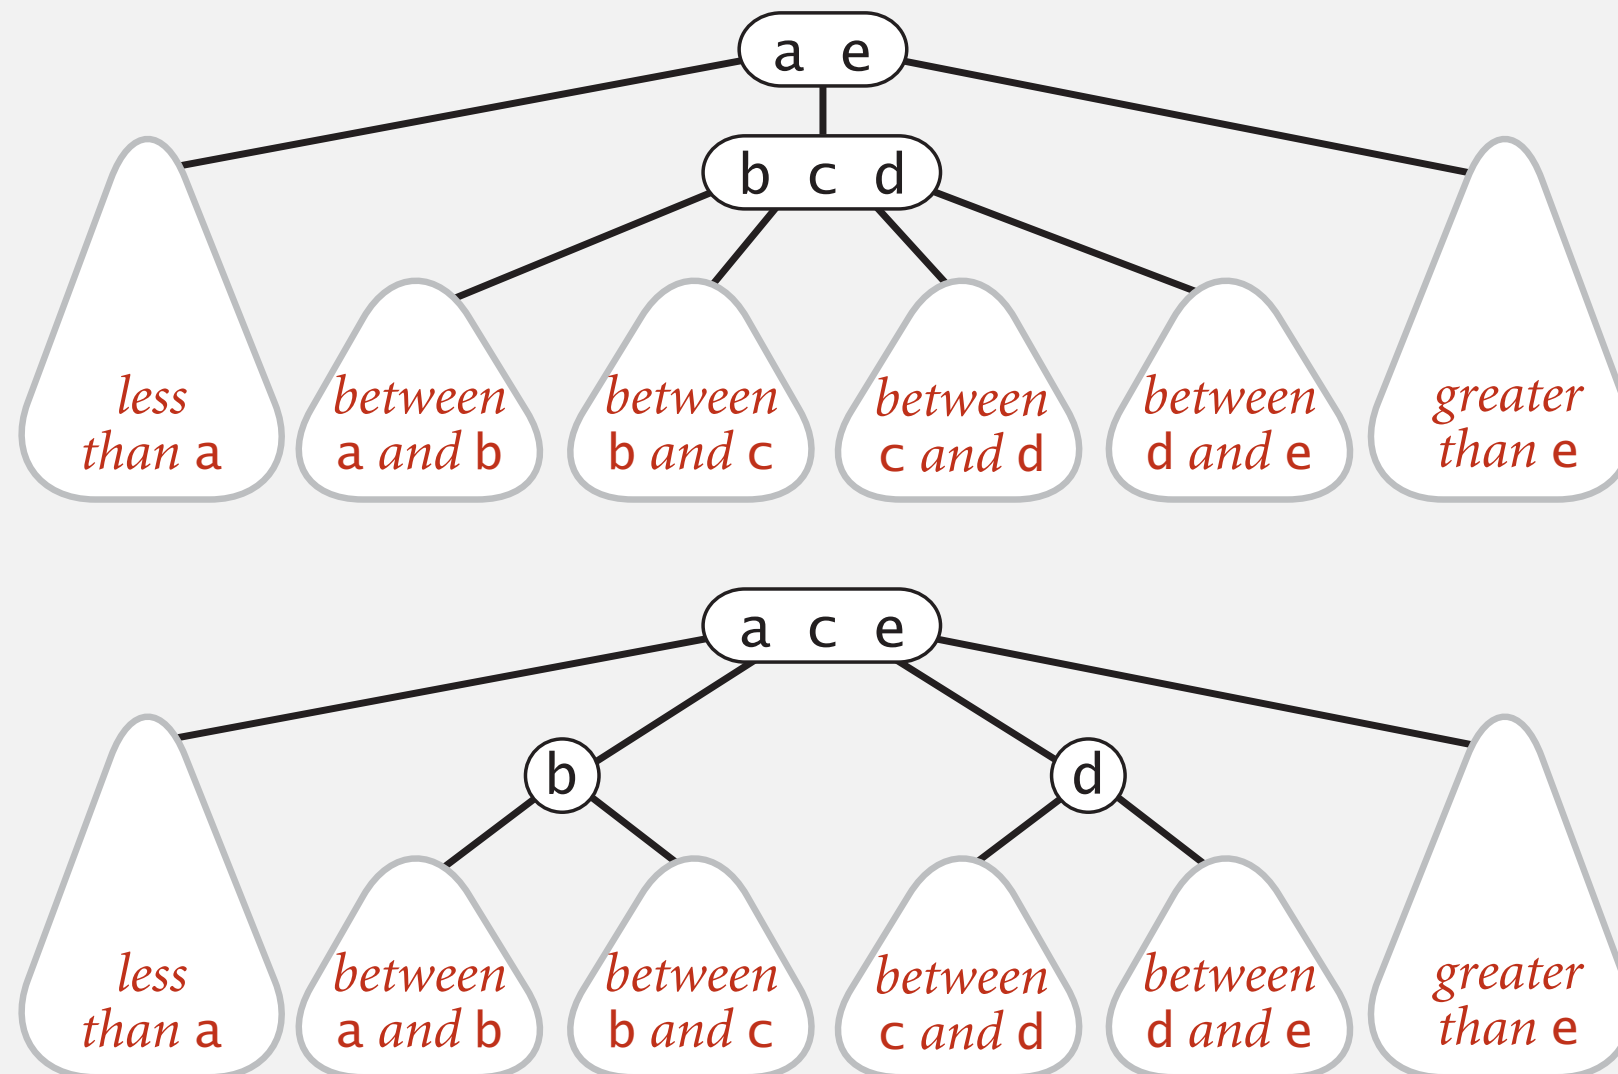


# Local transformations in a 2-3 tree

---

Splitting a 4-node is a **local** transformation: constant number of operations.

Splitting does not change the balance of the tree





# Global properties in a 2-3 tree

**Invariants.** Maintains symmetric order and perfect balance.

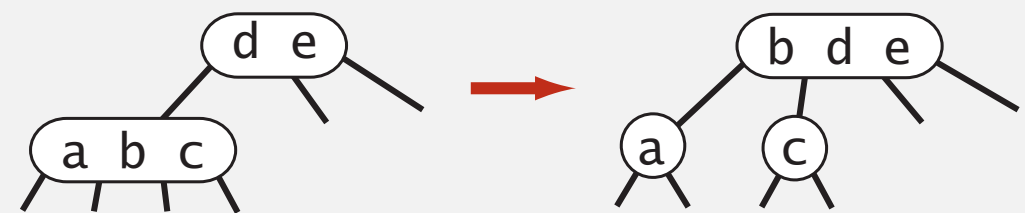
**Pf.** Each transformation maintains symmetric order and perfect balance.

root



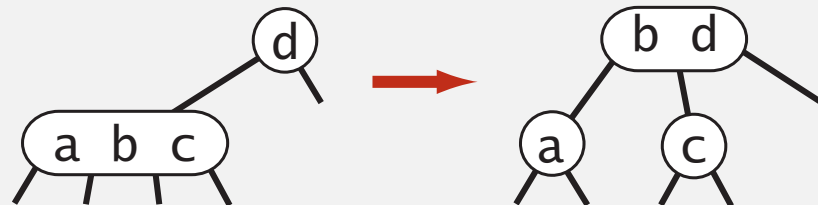
parent is a 3-node

left

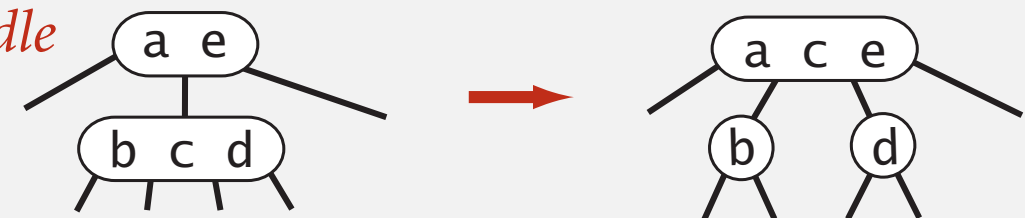


parent is a 2-node

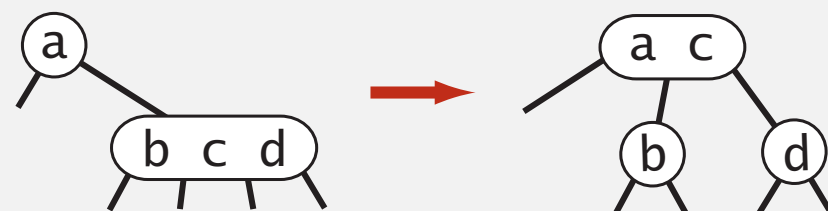
left



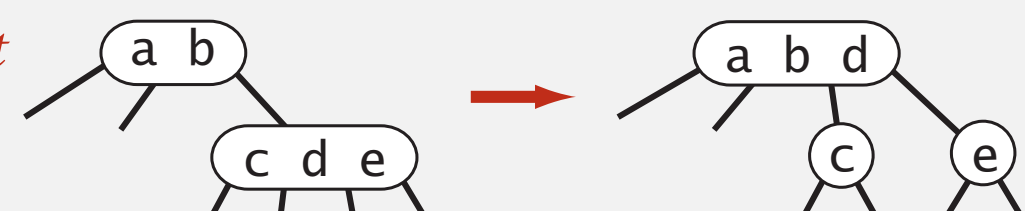
middle



right



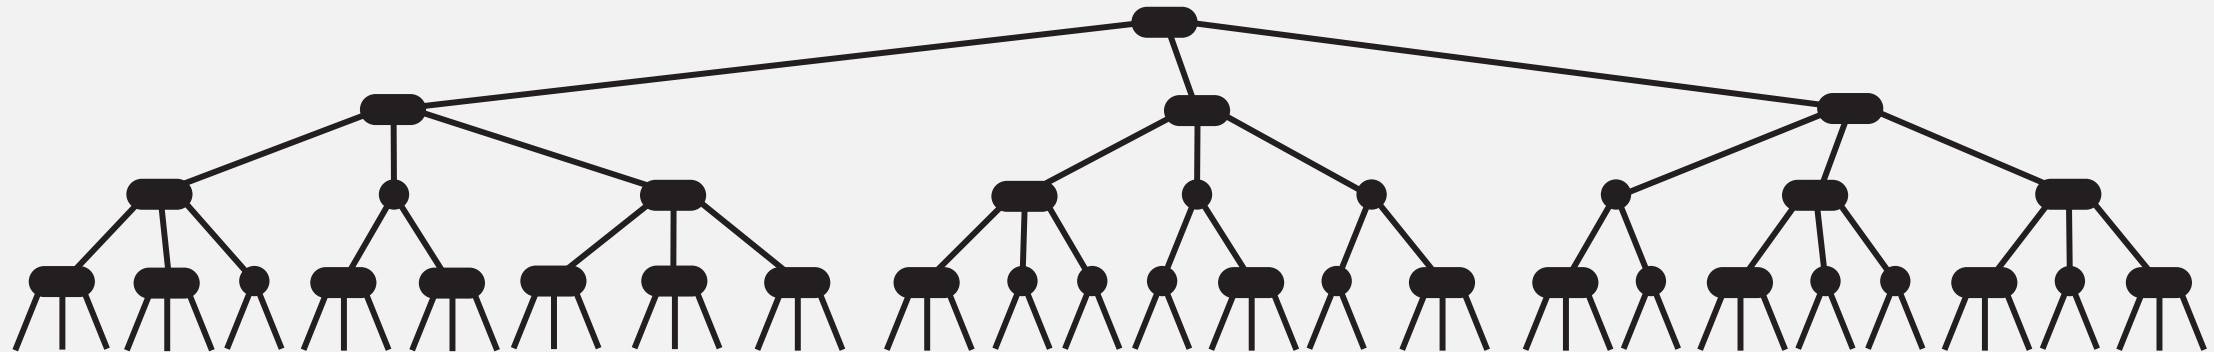
right



## 2-3 tree: performance

---

Perfect balance. Every path from root to null link has same length.



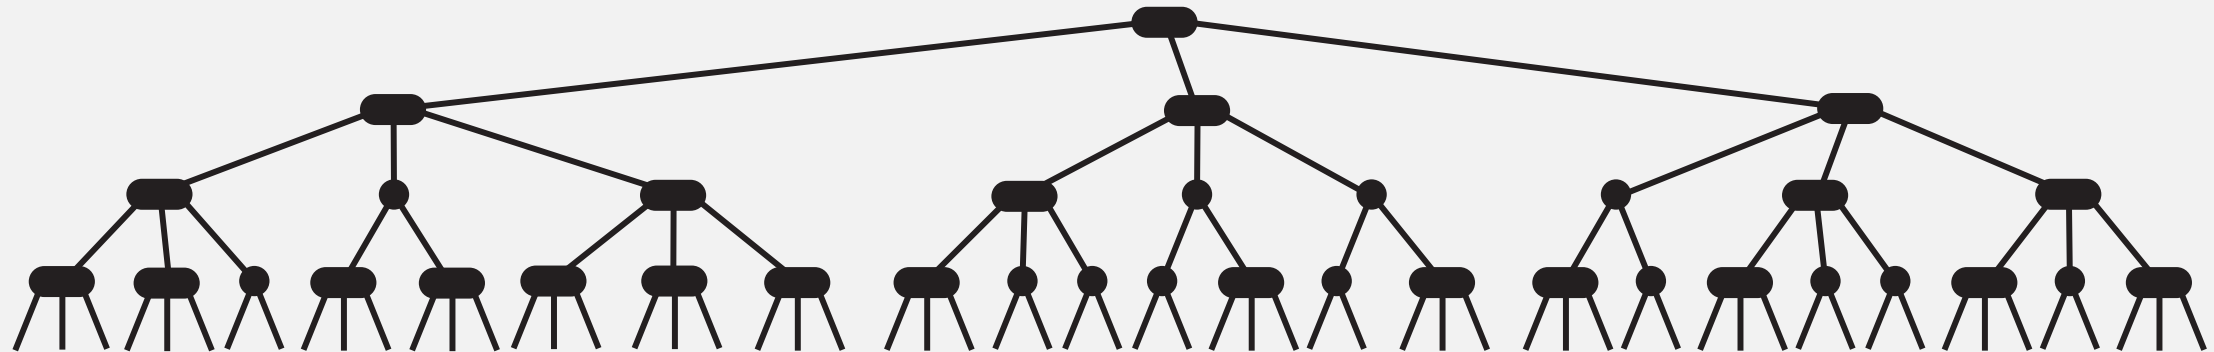
Tree height.

- Worst case:
- Best case:

## 2-3 tree: performance

---

**Perfect balance.** Every path from root to null link has same length.



### Tree height.

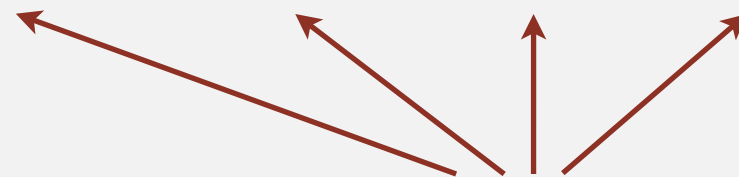
- Worst case:  $\log_2 N$ . [all 2-nodes]
- Best case:  $\log_3 N \approx .631 \log_2 N$ . [all 3-nodes]
- Between 12 and 20 for a million nodes.
- Between 18 and 30 for a billion nodes.

**Bottom line.** Guaranteed **logarithmic** performance for search and insert.

# ST implementations: summary

---

implementation	guarantee			average case		
	search	insert	delete	search hit	insert	delete
BST	$N$	$N$	$N$	$C \lg N$	$c \lg N$	$c \lg N$
2-3 tree	$c \lg N$	$c \lg N$	$c \lg N$	$c \lg N$	$c \lg N$	$c \lg N$



constant  $c$  depend upon implementation

# 2-3 TREES

Implementation

## 2-3 tree: implementation?

---

Direct implementation is complicated, because:

- Maintaining multiple node types is cumbersome.
- Need multiple compares to move down tree.
- Need to move back up the tree to split 4-nodes.
- Large number of cases for splitting.

*“ Beautiful algorithms are not always the most useful. ”*

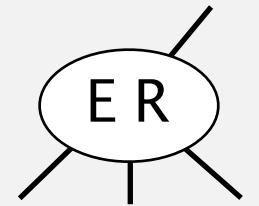
*— Donald Knuth*

**Bottom line.** Could do it, but there's a better way.

# How to implement 2-3 trees with binary trees?

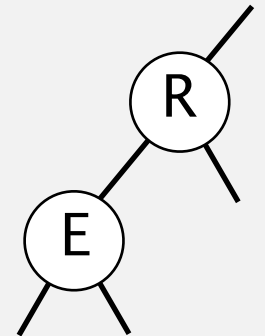
---

**Challenge.** How to represent a 3 node?



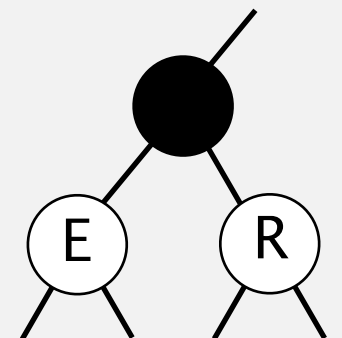
**Approach 1: regular BST.**

- No way to tell a 3-node from a 2-node.
- Cannot map from BST back to 2-3 tree.



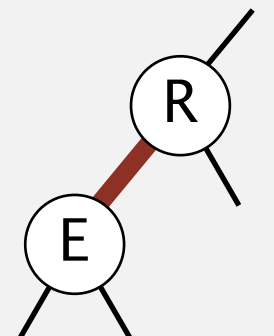
**Approach 2: regular BST with "glue" nodes.**

- Wastes space, wasted link.
- Code probably messy.



**Approach 3: regular BST with red "glue" links.**

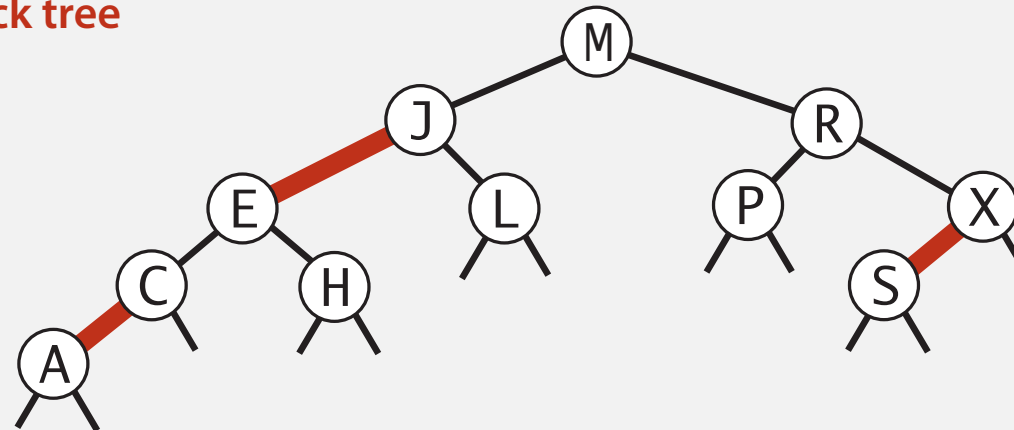
- Widely used in practice.
- Arbitrary restriction: red links lean left.



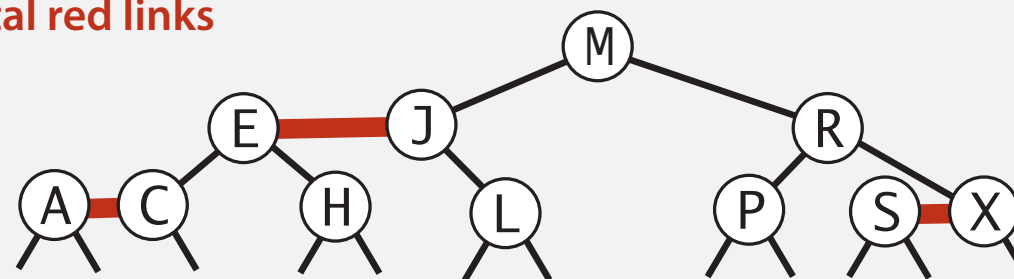
# Left-leaning red-black BSTs: 1-1 correspondence with 2-3 trees

**Key property.** 1–1 correspondence between 2–3 and LLRB.

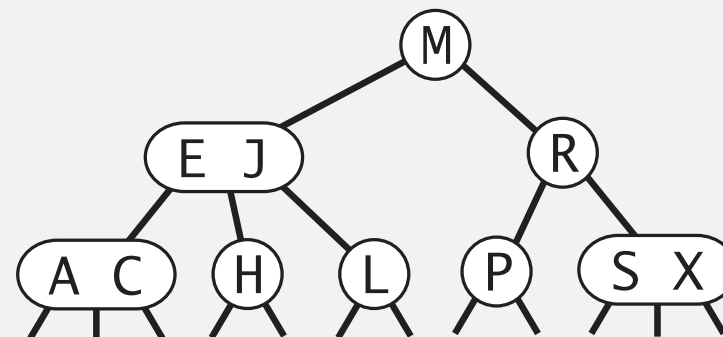
red-black tree



horizontal red links



2-3 tree





# Balanced trees in the wild

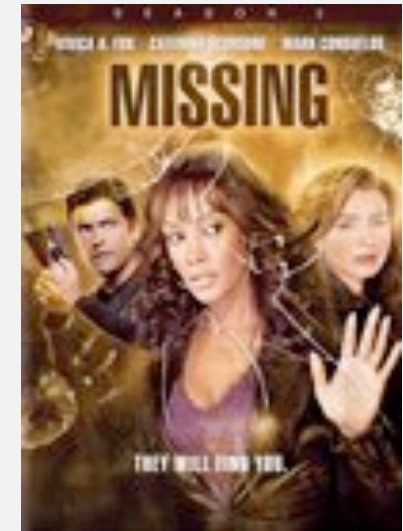
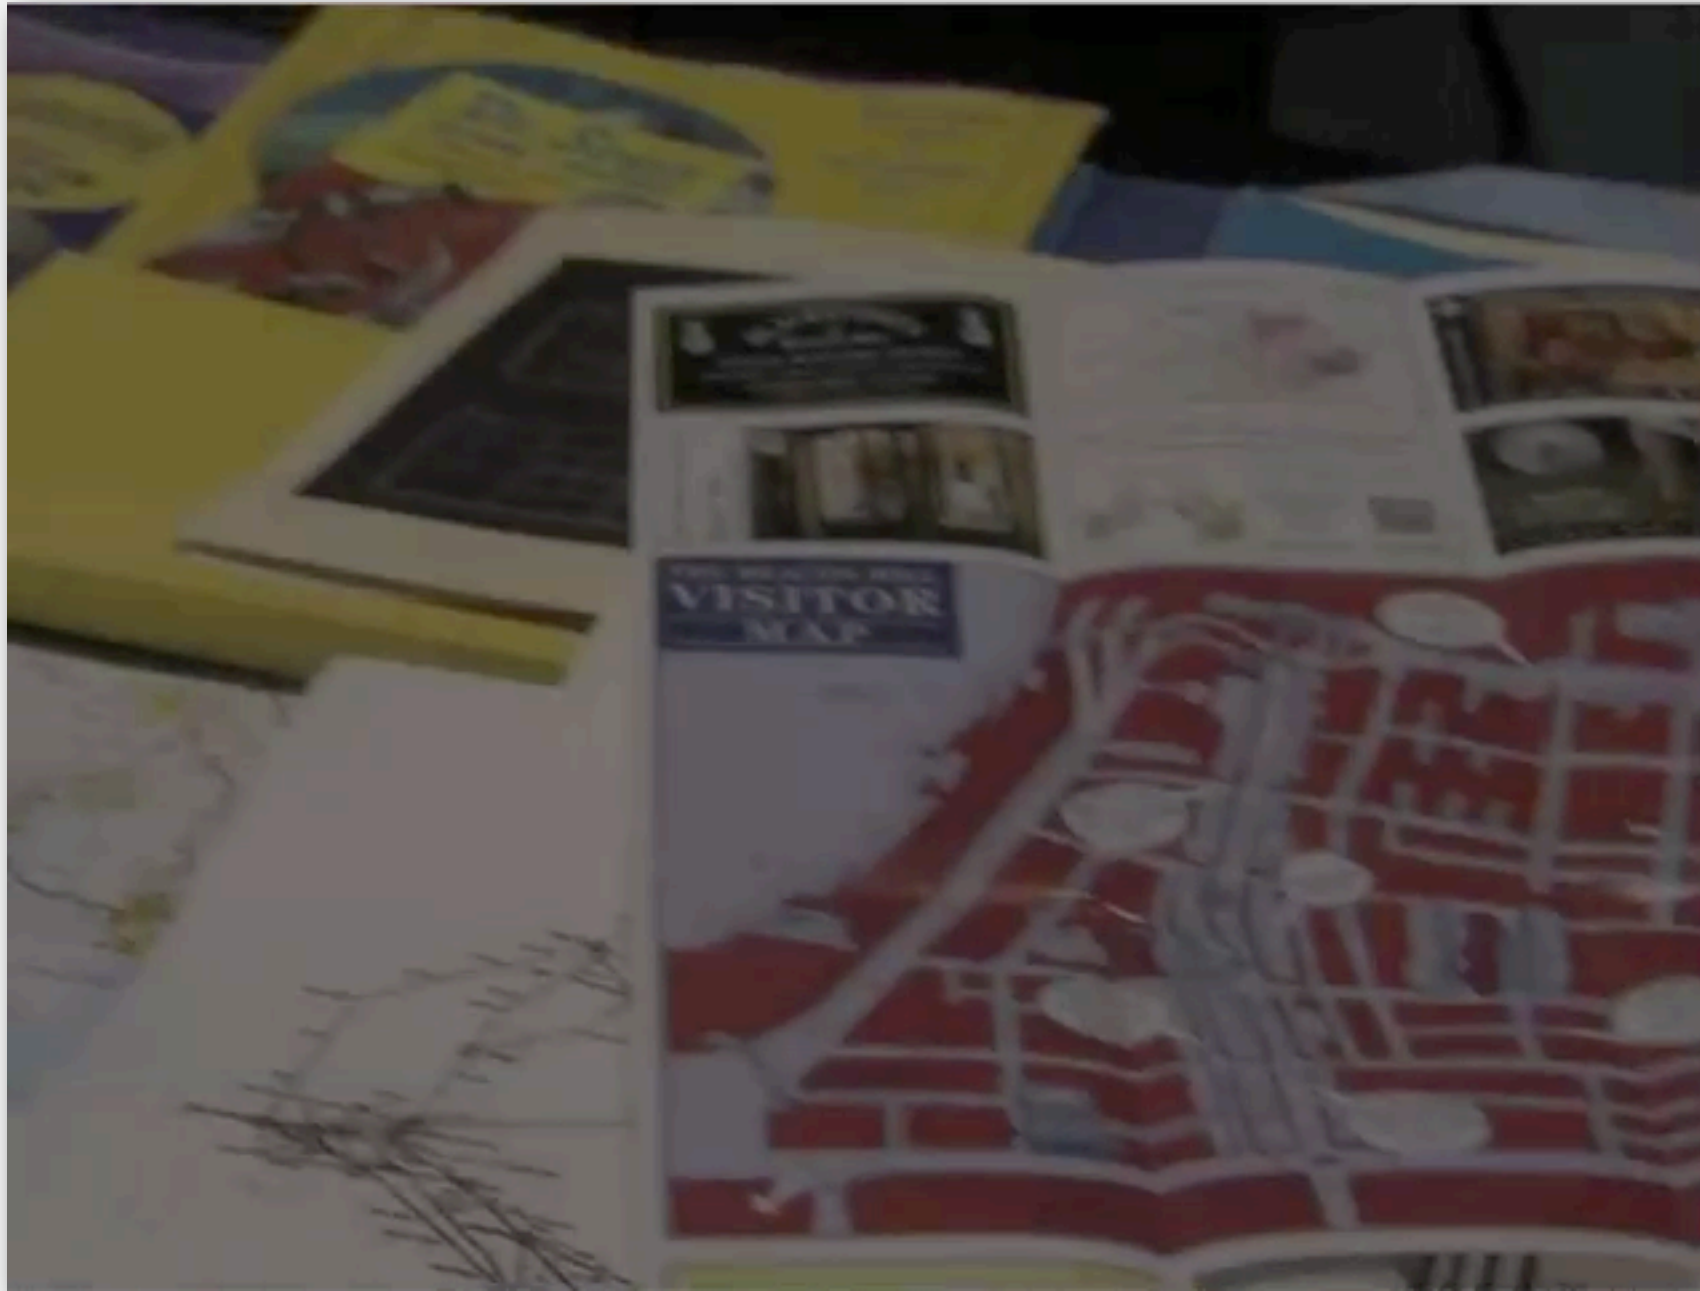
---

Red-black trees are widely used:

- Java: `java.util.TreeMap`, `java.util.TreeSet`.
- Linux kernel: completely fair scheduler, `linux/rbtree.h`.

# Red-black BSTs in the wild

---



# Red-black BSTs in the wild

---

## ACT FOUR

FADE IN:

48 INT. FBI HQ - NIGHT

48

Antonio is at THE COMPUTER as Jess explains herself to Nicole and Pollock. The CONFERENCE TABLE is covered with OPEN REFERENCE BOOKS, TOURIST GUIDES, MAPS and REAMS OF PRINTOUTS.

JESS

It was the red door again.

POLLOCK

I thought the red door was the storage container.

JESS

But it wasn't red anymore. It was black.

ANTONIO

So red turning to black means... what?

POLLOCK

Budget deficits? Red ink, black ink?

NICOLE

Yes. I'm sure that's what it is. But maybe we should come up with a couple other options, just in case.

Antonio refers to his COMPUTER SCREEN, which is filled with mathematical equations.

ANTONIO

It could be an algorithm from a binary search tree. A red-black tree tracks every simple path from a node to a descendant leaf with the same number of black nodes.

JESS

Does that help you with girls?

# War story: why red-black?

---

## Xerox PARC innovations. [1970s]

- Alto.
- GUI.
- Ethernet.
- Smalltalk.
- InterPress.
- Laser printing.
- Bitmapped display.
- WYSIWYG text editor.
- ...



**Xerox Alto**

### A DICHROMATIC FRAMEWORK FOR BALANCED TREES

Leo J. Guibas  
*Xerox Palo Alto Research Center,*  
Palo Alto, California, and  
*Carnegie-Mellon University*

and

Robert Sedgewick\*  
Program in Computer Science  
*Brown University*  
Providence, R. I.

#### ABSTRACT

In this paper we present a uniform framework for the implementation and study of balanced tree algorithms. We show how to imbed in this

the way down towards a leaf. As we will see, this has a number of significant advantages over the older methods. We shall examine a number of variations on a common theme and exhibit full implementations which are notable for their brevity. One implementation is examined carefully, and some properties about its

# War story: red-black BSTs

---

Telephone company contracted with database provider to build real-time database to store customer information.

## Database implementation.

- Red-black BST search and insert; (no rebalancing on deletion).
- Exceeding height limit of 80 triggered error-recovery process.

allows for up to  $2^{40}$  keys well over a trillion keys

## Extended telephone service outage.

- Main cause = height bounded exceeded!
- Telephone company sues database provider.
- Legal testimony:

*“ If implemented properly, the height of a red-black BST with  $N$  keys is at most  $2 \lg N$ . ” — expert witness*



B - TREES

## B-tree variants.

B-tree, Nodes contains keys and data.

B+ tree, - Data only stored in leafs (Deletion/Insertion simple just leaves)

B-trees (and variants) are widely used for file systems and databases.

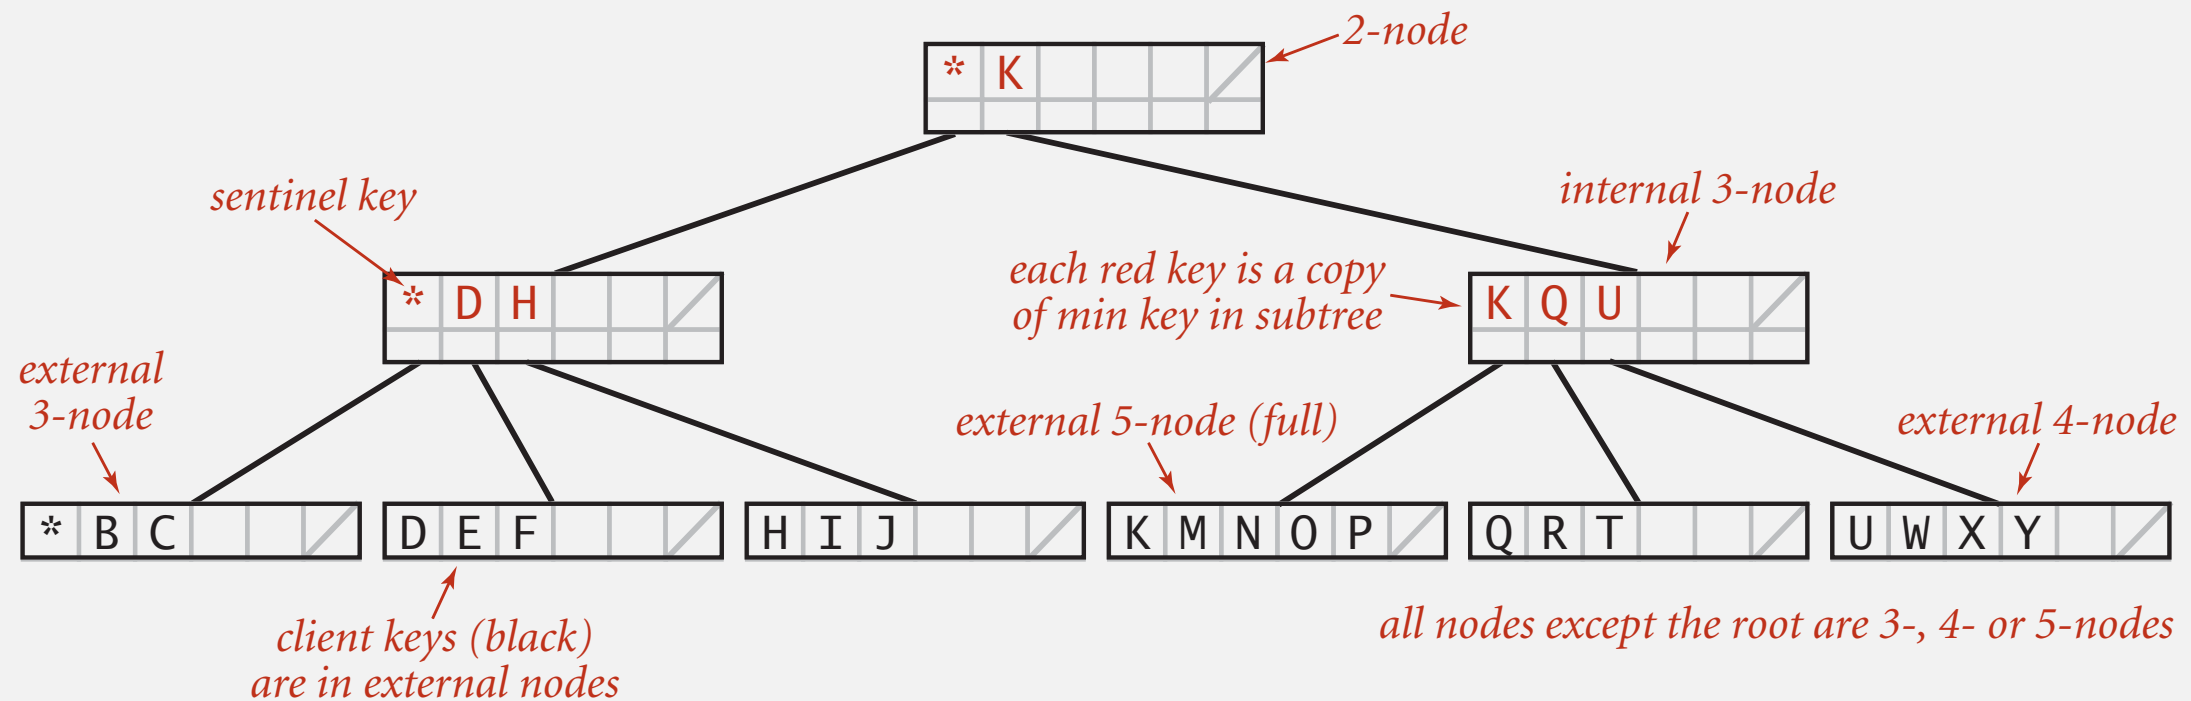
- Windows: NTFS.
- Mac: HFS, HFS+.
- Linux: ReiserFS, XFS, Ext3FS, JFS.
- Databases: ORACLE, DB2, INGRES, SQL, PostgreSQL.

# B-trees (Bayer-McCreight, 1972)

**B-tree.** Generalize 2-3 trees by allowing up to  $M - 1$  key per node.

- At least 2 key-link pairs at root.
- At least  $M / 2$  key-link pairs in other nodes.
- External nodes contain client keys.
- Internal nodes contain copies of keys to guide search.

choose  $M$  as large as possible so  
that  $M$  links fit in a page, e.g.,  $M = 1024$



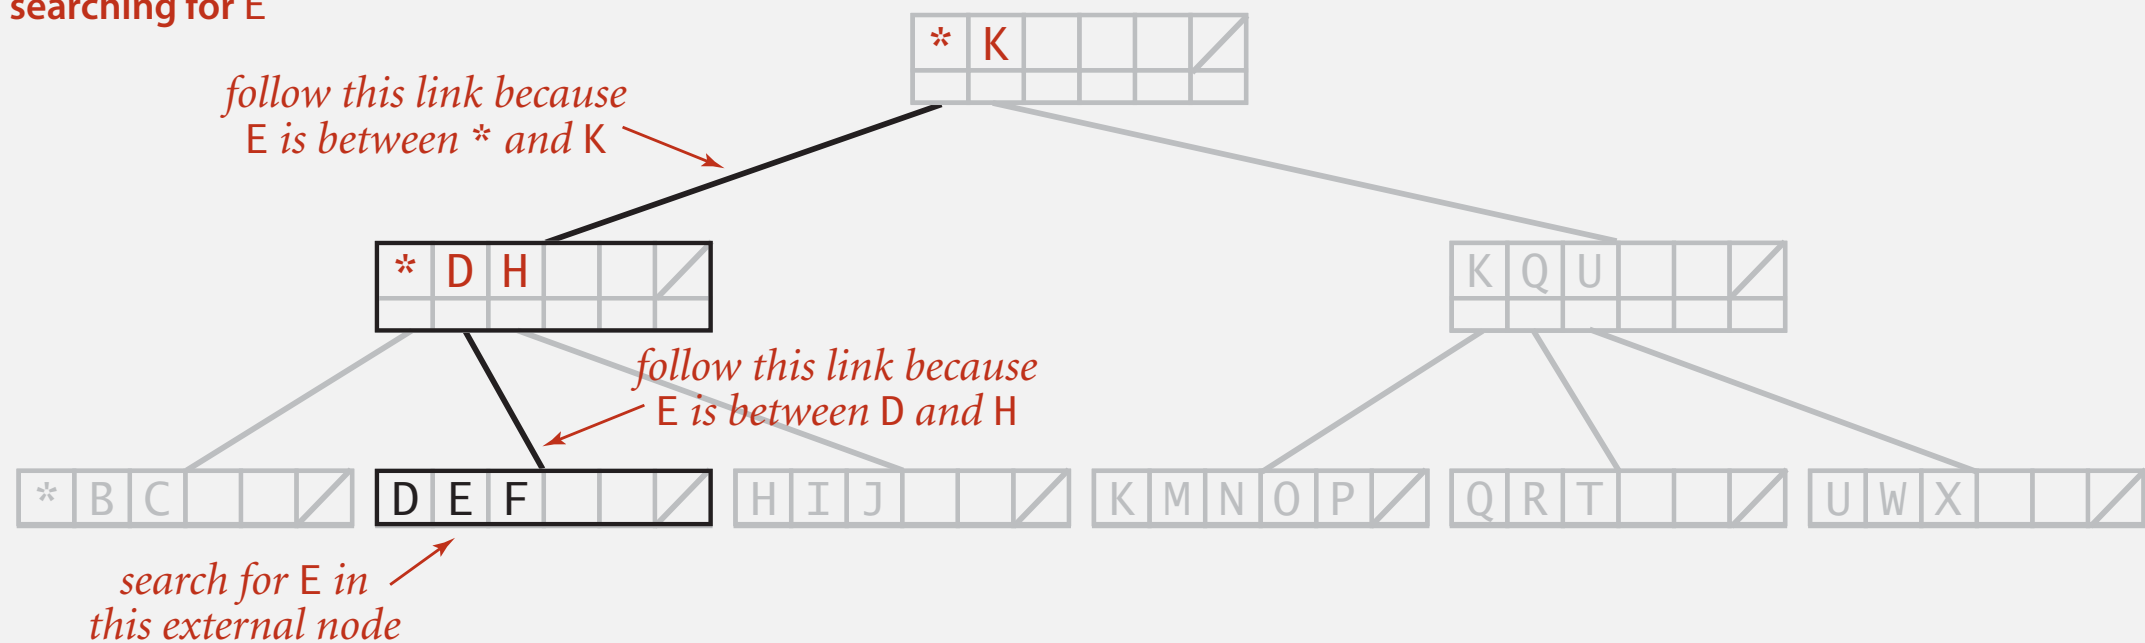
Anatomy of a B-tree set ( $M = 6$ )



# Searching in a B-tree

- Start at root.
- Find interval for search key and take corresponding link.
- Search terminates in external node.

searching for E

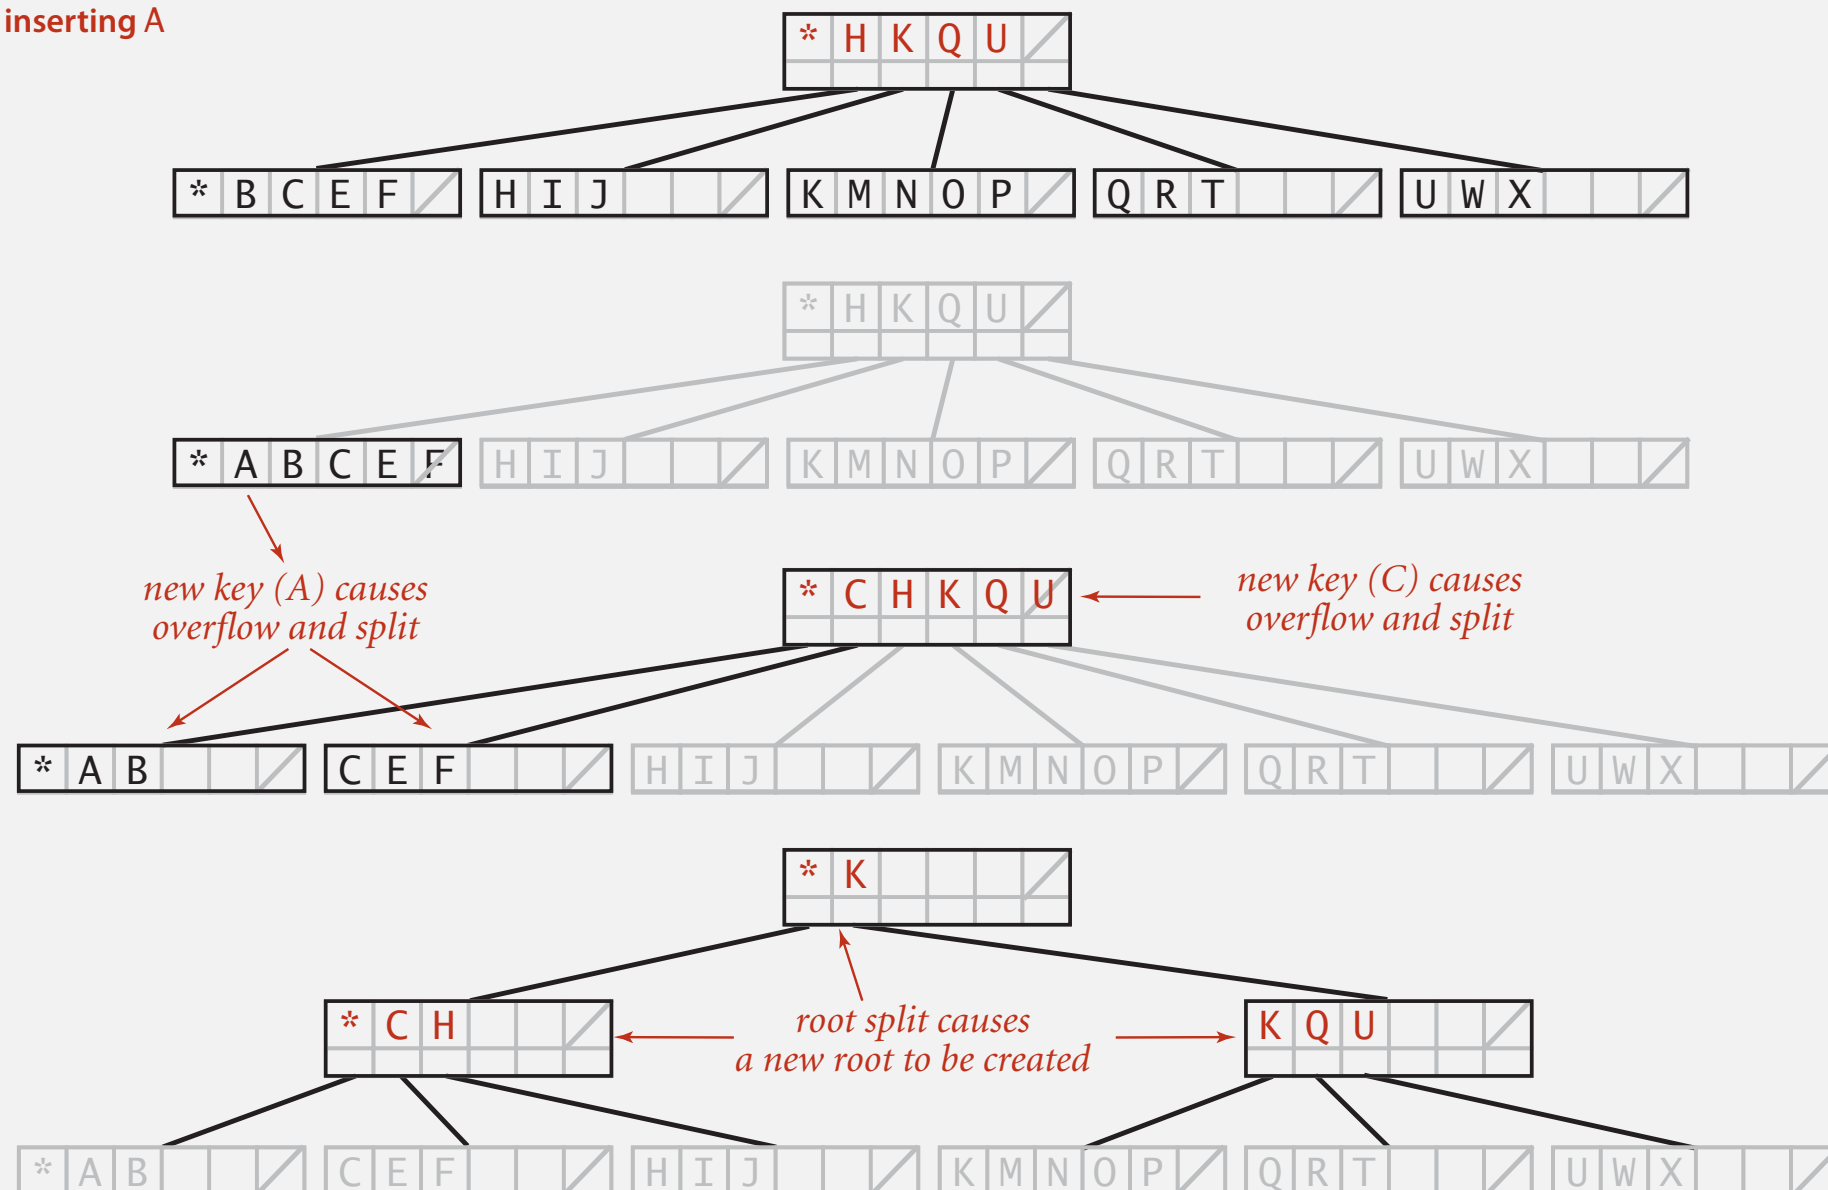


Searching in a B-tree set ( $M = 6$ )

# Insertion in a B-tree

- Search for new key.
- Insert at bottom.
- Split nodes with  $M$  key-link pairs on the way up the tree.

inserting A



Inserting a new key into a B-tree set

---



<http://algs4.cs.princeton.edu>

## 3.3 BALANCED SEARCH

---

- *2-3 search trees*
- *red-black BSTs*
- *B-trees*